

Statistical Learning for Biological and Information Systems

Problem Set #1

Chenye Yang cy2540@columbia.edu

September 25, 2019

P1

(a)

Ans:

$$\begin{aligned}(n-1)S^2 + n\bar{X}^2 &= \sum_{i=1}^n (X_i - \bar{X})^2 + n\bar{X}^2 \\ &= \sum_{i=1}^n X_i^2 - 2\bar{X} \sum_{i=1}^n X_i + 2n\bar{X}^2 \\ &= \sum_{i=1}^n X_i^2 - 2\bar{X}n\bar{X} + 2n\bar{X}^2 \\ &= \sum_{i=1}^n X_i^2\end{aligned}\tag{1}$$

(b)

Ans:

Because X_1, X_2, \dots, X_n are independent and identically distributed, using a as the expected value of the distribution, we have $\mathbb{E}(X_i) = \mathbb{E}(\bar{X}) = a$.

Thus:

$$\begin{aligned}
 \sum_{i=1}^n (X_i - \bar{X})^2 &= \sum_{i=1}^n [(X_i - a) - (\bar{X} - a)]^2 \\
 &= \sum_{i=1}^n [(X_i - a)^2 - 2(X_i - a)(\bar{X} - a) + (\bar{X} - a)^2] \\
 &= \sum_{i=1}^n (X_i - a)^2 - 2(\bar{X} - a) \left(\sum_{i=1}^n X_i - \sum_{i=1}^n a \right) + n(\bar{X} - a)^2 \\
 &= \sum_{i=1}^n (X_i - a)^2 - 2(\bar{X} - a)(n\bar{X} - na) + n(\bar{X} - a)^2 \\
 &= \sum_{i=1}^n (X_i - a)^2 - n(\bar{X} - a)^2
 \end{aligned} \tag{2}$$

Because:

$$\begin{aligned}
 \mathbb{E}[X_i - \mathbb{E}(X_i)]^2 &= \text{Var}(X_i) = \sigma^2 \\
 \mathbb{E}[\bar{X} - \mathbb{E}(\bar{X})]^2 &= \text{Var}(\bar{X}) = \sum_{i=1}^n \text{Var}(X_i)/n^2 = \sigma^2/n
 \end{aligned} \tag{3}$$

Thus:

$$\begin{aligned}
 \mathbb{E}S^2 &= \mathbb{E}\left[\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2\right] = \frac{1}{n-1} \mathbb{E}\left[\sum_{i=1}^n (X_i - a)^2 - n(\bar{X} - a)^2\right] \\
 &= \frac{1}{n-1} (n\sigma^2 - \sigma^2) = \sigma^2
 \end{aligned} \tag{4}$$

Namely, the S^2 is an unbiased estimator of σ^2

(c)

Ans:

Because $X_i \sim N(\mu, \sigma^2)$ ($i = 1, 2, \dots, n$) and X_1, X_2, \dots, X_n are i.i.d.

Therefore, \bar{X} and $X_i - \bar{X}$ ($i = 1, 2, \dots, n$) have normal/Gaussian distribution, and we have:

$$\begin{aligned}
 \mathbb{E}(\bar{X}) &= \mathbb{E}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n \mathbb{E}(X_i) = \mu \\
 \text{Var}(\bar{X}) &= \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{1}{n} \sigma^2 \\
 \mathbb{E}(X_i - \bar{X}) &= \mathbb{E}(X_i) - \mathbb{E}(\bar{X}) = \mu - \mu = 0 \\
 \text{Var}(X_i - \bar{X}) &= \text{Var}(X_i) - \text{Var}(\bar{X}) = \sigma^2 - \frac{1}{n} \sigma^2 = \frac{n-1}{n} \sigma^2
 \end{aligned} \tag{5}$$

Thus:

$$\begin{aligned}
 \text{Cov}(\bar{X}, X_i - \bar{X}) &= \mathbb{E}[(\bar{X} - \mathbb{E}(\bar{X}))((X_i - \bar{X}) - \mathbb{E}(X_i - \bar{X}))] \\
 &= \mathbb{E}[(\bar{X} - \mu)(X_i - \bar{X})] \\
 &= \mathbb{E}[\bar{X}X_i - \bar{X}^2 - \mu X_i + \mu\bar{X}] \\
 &= \mathbb{E}(\bar{X}X_i) - \mathbb{E}(\bar{X}^2) - \mu^2 + \mu^2 \\
 &= \mathbb{E}(\bar{X}X_i) - [\text{Var}(\bar{X}) + \mu^2] \\
 &= \mathbb{E}(\bar{X}X_i) - \left(\frac{\sigma^2}{n} + \mu^2\right)
 \end{aligned} \tag{6}$$

Because:

$$\begin{aligned}
 \mathbb{E}(\bar{X}X_i) &= \mathbb{E}\left(X_i \frac{1}{n} \sum_{j=1}^n X_j\right) = \frac{1}{n} \mathbb{E}\left(\sum_{j=1}^n X_j X_i\right) \\
 &= \frac{1}{n} \mathbb{E}\left[\sum_{j=1, j \neq i}^n (X_j X_i) + X_i^2\right] \\
 &= \frac{1}{n} \left[\sum_{j=1, j \neq i}^n \mathbb{E}(X_j)\mathbb{E}(X_i) + \mathbb{E}(X_i^2)\right] \\
 &= \frac{1}{n} [(n-1)\mu^2 + \text{Var}(X_i) + \mu^2] \\
 &= \mu^2 + \frac{1}{n} \sigma^2
 \end{aligned} \tag{7}$$

Therefore:

$$\begin{aligned}
 \text{Cov}(\bar{X}, X_i - \bar{X}) &= (\mu^2 + \frac{1}{n} \sigma^2) - \left(\frac{\sigma^2}{n} + \mu^2\right) = 0 \\
 \text{Corr}(\bar{X}, X_i - \bar{X}) &= \text{Cov}(\bar{X}, X_i - \bar{X}) / \left(\sqrt{\frac{1}{n} \sigma^2} \sqrt{\frac{n-1}{n} \sigma^2}\right) = 0
 \end{aligned} \tag{8}$$

Considering \bar{X} and $X_i - \bar{X}$ ($i = 1, 2, \dots, n$) have normal/Gaussian distribution, thus \bar{X} is independent of $X_i - \bar{X}$ ($i = 1, 2, \dots, n$).

(d)

Ans:

Let $Z_i = \frac{X_i - \mu}{\sigma}$ $i = 1, 2, \dots, n$, thus Z_1, Z_2, \dots, Z_n are independent and $Z_i \sim N(0, 1)$ $i = 1, 2, \dots, n$. Also:

$$\begin{aligned}
 \bar{Z} &= \frac{1}{n} \sum_{i=1}^n Z_i = \frac{1}{n} \sum_{i=1}^n \frac{X_i - \mu}{\sigma} = \frac{\bar{X} - \mu}{\sigma} \\
 \frac{(n-1)S^2}{\sigma^2} &= \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{\sigma^2} = \sum_{i=1}^n \left[\frac{(X_i - \mu) - (\bar{X} - \mu)}{\sigma} \right]^2 \\
 &= \sum_{i=1}^n (Z_i - \bar{Z})^2 = \sum_{i=1}^n (Z_i^2 - 2Z_i\bar{Z} + \bar{Z}^2) \\
 &= \sum_{i=1}^n Z_i^2 - 2\bar{Z} \sum_{i=1}^n Z_i + n\bar{Z}^2 = \sum_{i=1}^n Z_i^2 - n\bar{Z}^2
 \end{aligned} \tag{9}$$

Choose a matrix $\mathbf{A} = (a_{ij})$, in which all the entries in the first row are $1/\sqrt{n}$, and each row vector of the matrix should be orthogonal. Do:

$$\mathbf{Y} = \mathbf{AZ}$$

In which

$$\mathbf{Y} = (Y_1, Y_2, \dots, Y_n)^T$$

$$\mathbf{Z} = (Z_1, Z_2, \dots, Z_n)^T$$

$$\mathbf{A} = (\mathbf{a}_1^T, \mathbf{a}_2^T, \dots, \mathbf{a}_n^T)^T$$

Because Y_i is the linear combination of Z_i and $Z_i \sim N(0, 1)$ $i = 1, 2, \dots, n$, therefore Y_i $i = 1, 2, \dots, n$ have normal/Gaussian distribution and:

$$\begin{aligned} \mathbb{E}(Y_i) &= \mathbb{E}\left(\sum_{j=1}^n a_{ij} Z_j\right) = \sum_{j=1}^n a_{ij} \mathbb{E}(Z_j) = 0 \\ \text{Var}(Y_i) &= \text{Var}\left(\sum_{j=1}^n a_{ij} Z_j\right) = \sum_{j=1}^n a_{ij}^2 = \langle \mathbf{a}_i, \mathbf{a}_i \rangle = 1 \end{aligned} \tag{10}$$

Because $\text{Cov}(Z_k, Z_h)$ equals 1 when $k = h$ and equals 0 when $k \neq h$:

$$\begin{aligned} \text{Cov}(Y_i, Y_j) &= \text{Cov}\left(\sum_{k=1}^n a_{ik} Z_k, \sum_{h=1}^n a_{jh} Z_h\right) \\ &= \sum_{k=1}^n \sum_{h=1}^n a_{ik} a_{jh} \text{Cov}(Z_k, Z_h) \\ &= \sum_{k=1}^n a_{ik} a_{jk} \\ &= \langle \mathbf{a}_i, \mathbf{a}_j \rangle \\ &= \begin{cases} 1 & , i = j \\ 0 & , i \neq j \end{cases} \\ &= \text{Cov}(Z_i, Z_j) \end{aligned} \tag{11}$$

Therefore, Y_i $i = 1, 2, \dots, n$ has no correlation with each other. Considering $Y_i \sim N(0, 1)$, thus Y_1, Y_2, \dots, Y_n are independent. Also:

$$\begin{aligned} Y_1 &= \sum_{j=1}^n a_{1j} Z_j = \sum_{j=1}^n \frac{1}{\sqrt{n}} Z_j = \sqrt{n} \bar{Z} \\ \sum_{i=1}^n Y_i^2 &= \mathbf{Y}^T \mathbf{Y} = (\mathbf{AZ})^T (\mathbf{AZ}) = \mathbf{Z}^T (\mathbf{A} \mathbf{A}^T)^T \mathbf{Z} = \mathbf{Z}^T \mathbf{I} \mathbf{Z} = \sum_{i=1}^n Z_i^2 \end{aligned} \tag{12}$$

Therefore:

$$\frac{(n-1)S^2}{\sigma^2} = \sum_{i=1}^n Z_i^2 - n\bar{Z}^2 = \sum_{i=1}^n (Y_i^2) - Y_1^2 = \sum_{i=2}^n (Y_i^2) \tag{13}$$

Because $\bar{X} = \sigma \bar{Z} + \mu = \frac{\sigma}{\sqrt{n}} Y_1 + \mu$ only depends on Y_1 and $S^2 = \frac{\sigma^2}{n-1} \sum_{i=2}^n (Y_i^2)$ only depends on Y_2, Y_3, \dots, Y_n , also Y_1, Y_2, \dots, Y_n are independent, therefore, the sample mean \bar{X} is independent of the sample variance S^2 .

P2

Ans:

For simple linear regression, we have $Y \approx \beta_0 + \beta_1 X$, in which the value of β_0 and β_1 to minimize the Residual Sum of Squares (RSS) is as follow:

$$\begin{aligned}\hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{\text{Cov}(X, Y)}{\text{Var}(X)} \\ \hat{\beta}_0 &= \bar{y} - \hat{\beta}_1 \bar{x}\end{aligned}\tag{14}$$

Considering:

$$\begin{aligned}\sum_{i=1}^n (y_i - \hat{y}_i)^2 &= \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2 \\ &= \sum_{i=1}^n (y_i - \bar{y} + \hat{\beta}_1 \bar{x} - \hat{\beta}_1 x_i)^2 \\ &= \sum_{i=1}^n [y_i - \bar{y} - \hat{\beta}_1 (x_i - \bar{x})]^2 \\ &= \sum_{i=1}^n [(y_i - \bar{y})^2 - 2\hat{\beta}_1 (y_i - \bar{y})(x_i - \bar{x}) + \hat{\beta}_1^2 (x_i - \bar{x})^2] \\ &= \sum_{i=1}^n (y_i - \bar{y})^2 - 2\hat{\beta}_1 \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x}) + \hat{\beta}_1^2 \sum_{i=1}^n (x_i - \bar{x})^2\end{aligned}\tag{15}$$

Thus:

$$\begin{aligned}R^2 &= 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \\ &= 1 - \frac{\sum_{i=1}^n (y_i - \bar{y})^2 - 2\hat{\beta}_1 \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x}) + \hat{\beta}_1^2 \sum_{i=1}^n (x_i - \bar{x})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \\ &= 1 - 1 + 2\hat{\beta}_1 \frac{\sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})}{\sum_{i=1}^n (y_i - \bar{y})^2} - \hat{\beta}_1^2 \frac{\sum_{i=1}^n (x_i - \bar{x})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \\ &= 2\hat{\beta}_1 \frac{\text{Cov}(X, Y)}{\text{Var}(Y)} - \hat{\beta}_1^2 \frac{\text{Var}(X)}{\text{Var}(Y)}\end{aligned}\tag{16}$$

Replace $\hat{\beta}_1$ in the equation:

$$\begin{aligned}R^2 &= 2 \frac{\text{Cov}(X, Y)}{\text{Var}(X)} \frac{\text{Cov}(X, Y)}{\text{Var}(Y)} - \left[\frac{\text{Cov}(X, Y)}{\text{Var}(X)} \right]^2 \frac{\text{Var}(X)}{\text{Var}(Y)} \\ &= \frac{\text{Cov}(X, Y)^2}{\text{Var}(X)\text{Var}(Y)} \\ &= r^2\end{aligned}\tag{17}$$

Therefore, in the case of simple linear regression of Y onto X , the R^2 statistic is equal to the square of the correlation coefficient between X and Y (r^2).

P3**(a)****Code:**

```

1 > set.seed(1)
2 > x=rnorm(100, mean=0, sd=1)

```

Result:

```

1 x
2  [1] -0.626453811  0.183643324 -0.835628612  1.595280802  0.329507772
3  [6] -0.820468384  0.487429052  0.738324705  0.575781352 -0.305388387
4  [11]  1.511781168  0.389843236 -0.621240581 -2.214699887  1.124930918
5  [16] -0.044933609 -0.016190263  0.943836211  0.821221195  0.593901321
6  [21]  0.918977372  0.782136301  0.074564983 -1.989351696  0.619825748
7  [26] -0.056128740 -0.155795507 -1.470752384 -0.478150055  0.417941560
8  [31]  1.358679552 -0.102787727  0.387671612 -0.053805041 -1.377059557
9  [36] -0.414994563 -0.394289954 -0.059313397  1.100025372  0.763175748
10 [41] -0.164523596 -0.253361680  0.696963375  0.556663199 -0.688755695
11 [46] -0.707495157  0.364581962  0.768532925 -0.112346212  0.881107726
12 [51]  0.398105880 -0.612026393  0.341119691 -1.129363096  1.433023702
13 [56]  1.980399899 -0.367221476 -1.044134626  0.569719627 -0.135054604
14 [61]  2.401617761 -0.039240003  0.689739362  0.028002159 -0.743273209
15 [66]  0.188792300 -1.804958629  1.465554862  0.153253338  2.172611670
16 [71]  0.475509529 -0.709946431  0.610726353 -0.934097632 -1.253633400
17 [76]  0.291446236 -0.443291873  0.001105352  0.074341324 -0.589520946
18 [81] -0.568668733 -0.135178615  1.178086997 -1.523566800  0.593946188
19 [86]  0.332950371  1.063099837 -0.304183924  0.370018810  0.267098791
20 [91] -0.542520031  1.207867806  1.160402616  0.700213650  1.586833455
21 [96]  0.558486426 -1.276592208 -0.573265414 -1.224612615 -0.473400636

```

(b)**Code:**

```

1 > eps=rnorm(100, mean=0, sd=0.25)

```

Result:

```

1 eps
2  [1] -0.155091669  0.010528968 -0.227730412  0.039507193 -0.163646161
3  [6]  0.441821817  0.179176869  0.227543557  0.096046339  0.420544020
4  [11] -0.158934113 -0.115411183  0.358070560 -0.162674088 -0.051845186
5  [16] -0.098201982 -0.079998217 -0.069778326  0.123547083 -0.044332621
6  [21] -0.126489366  0.335759706 -0.053644852 -0.044889133 -0.025047685
7  [26]  0.178166577 -0.018391101 -0.009408543 -0.170415120 -0.081067568
8  [31]  0.015040110 -0.147223622  0.132874048 -0.379598520  0.076639465

```

9	[36]	-0.384112456	-0.075244032	-0.132069976	-0.163023695	-0.014224194
10	[41]	-0.478589856	0.294145828	-0.416243109	-0.115882600	-0.278980026
11	[46]	-0.187704750	0.521791636	0.004348905	-0.321575133	-0.410151384
12	[51]	0.112546775	-0.004639958	-0.079517094	-0.232340537	-0.371865078
13	[56]	-0.268798074	0.250007201	-0.155316674	-0.346106712	0.467322656
14	[61]	0.106275094	-0.059661775	0.264620762	0.221605663	-0.154810762
15	[66]	0.551525616	-0.063756758	-0.356123663	-0.036099900	0.051884585
16	[71]	0.576994600	0.026450592	0.114249701	-0.019288234	-0.083500211
17	[76]	-0.008681507	0.196909901	0.518811252	0.256848110	0.301977100
18	[81]	-0.307830855	0.245973893	0.054981201	-0.366812507	0.130255686
19	[86]	-0.039688651	0.366146828	-0.191520500	-0.107552938	-0.231527374
20	[91]	-0.044275990	0.100502945	-0.182937043	0.207593292	-0.302020697
21	[96]	-0.261996103	0.360289427	-0.253961866	0.102993678	-0.095269013

(c)

Code:

```
1 > y = -1 + 0.5*x + eps
2 > length(y)
```

Result:

1	y
2	[1] -1.4683186 -0.8976494 -1.6455447 -0.1628524 -0.9988923 -0.9684124 -0.5771086
3	[8] -0.4032941 -0.6160630 -0.7321502 -0.4030435 -0.9204896 -0.9525497 -2.2700240
4	[15] -0.4893797 -1.1206688 -1.0880933 -0.5978602 -0.4658423 -0.7473820 -0.6670007
5	[22] -0.2731721 -1.0163624 -2.0395650 -0.7151348 -0.8498978 -1.0962889 -1.7447847
6	[29] -1.4094901 -0.8720968 -0.3056201 -1.1986175 -0.6732901 -1.4065010 -1.6118903
7	[36] -1.5916097 -1.2723890 -1.1617267 -0.6130110 -0.6326363 -1.5608517 -0.8325350
8	[43] -1.0677614 -0.8375510 -1.6233579 -1.5414523 -0.2959174 -0.6113846 -1.3777482
9	[50] -0.9695975 -0.6884003 -1.3106532 -0.9089572 -1.7970221 -0.6553532 -0.2785981
10	[57] -0.9336035 -1.6773840 -1.0612469 -0.6002046 0.3070840 -1.0792818 -0.3905096
11	[64] -0.7643933 -1.5264474 -0.3540782 -1.9662361 -0.6233462 -0.9594732 0.1381904
12	[71] -0.1852506 -1.3285226 -0.5803871 -1.4863370 -1.7103169 -0.8629584 -1.0247360
13	[78] -0.4806361 -0.7059812 -0.9927834 -1.5921652 -0.8216154 -0.3559753 -2.1285959


```

14 [85] -0.5727712 -0.8732135 -0.1023033 -1.3436125 -0.9225435
    -1.0979780 -1.3155360
15 [92] -0.2955632 -0.6027357 -0.4422999 -0.5086040 -0.9827529
    -1.2780067 -1.5405946
16 [99] -1.5093126 -1.3319693
17
18 length(y)
19 [1] 100

```

Ans:

The length of the vector y is 100.

$\beta_0 = -1$ and $\beta_1 = 0.5$

(d)


Code:

```

1 > plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")

```

Result:



Ans:

There is an approximate linear relationship between Y and X.

(e)

Code:

```

1 > lm.fit = lm(y~x)
2 > summary(lm.fit)
3 > plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
4 > abline(lm.fit)


```

Result:

```

1 Call:
2 lm(formula = y ~ x)
3
4 Residuals:
5      Min       1Q   Median       3Q      Max
6 -0.46921 -0.15344 -0.03487  0.13485  0.58654
7
8 Coefficients:
9             Estimate Std. Error t value Pr(>|t|)
10 (Intercept) -1.00942    0.02425  -41.63  <2e-16 ***
11 x           0.49973    0.02693   18.56  <2e-16 ***
12 ---
13 Signif. codes:  0      ***    0.001    **    0.01    *    0.05    .
14                  0.1      1
15 Residual standard error: 0.2407 on 98 degrees of freedom
16 Multiple R-squared:  0.7784,    Adjusted R-squared:  0.7762
17 F-statistic: 344.3 on 1 and 98 DF,  p-value: < 2.2e-16

```

**Ans:**

$$\hat{\beta}_0 = -1.00942 \quad \hat{\beta}_1 = 0.49973 \quad \beta_0 = -1 \quad \beta_1 = 0.5$$

With the $\Pr(> |t|) < 2e - 16$, we will reject H_0 and say $\hat{\beta}_0$ and $\hat{\beta}_1$ are not 0.

With p-value $< 2.2e - 16$ and $R^2 = 0.7784$, we can say there exist relationship between Y and X and our linear model fits Y and X well.

The estimated value of β_0 and β_1 are very close to their true value.

(f)

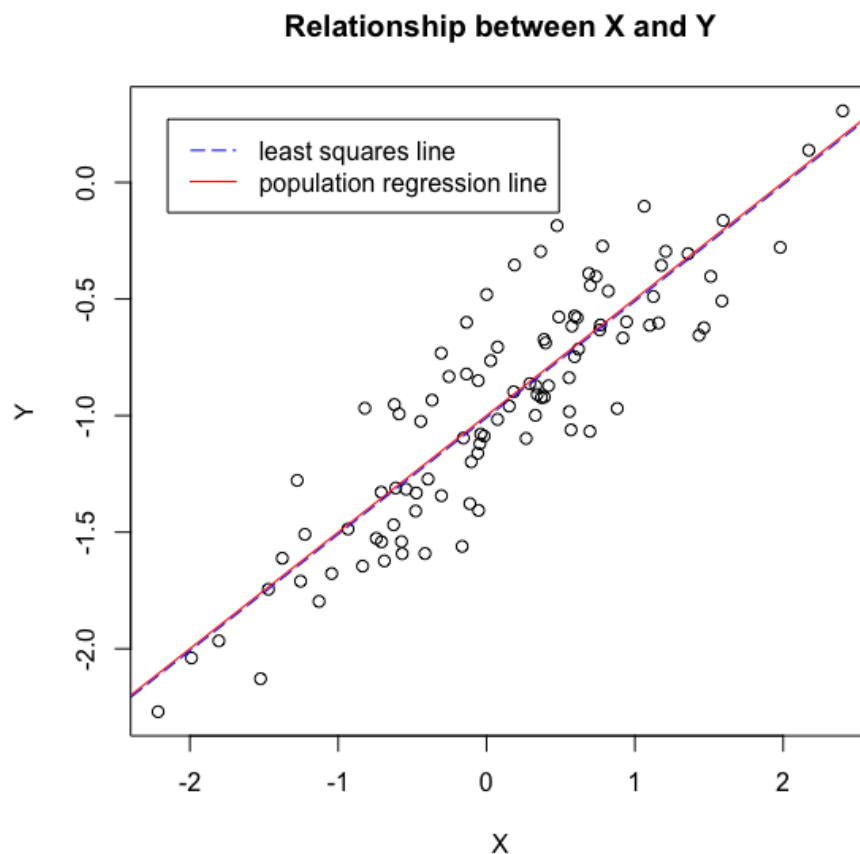
Code:

```

1 > par(col='black')
2 > plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
3 > abline(lm.fit, col='blue', lty=5)
4 > abline(a=-1, b=0.5, col='red', lty=1)
5 > legend('topleft', inset=0.05, c('least_squares_line', 'population_
    regression_line'), lty=c(5, 1), col=c('blue', 'red'), bty = "o")

```

Result:



(g)

Code:

```

1 # polynomial regression
2 Poly_fit = lm(y ~ poly(x,2))
3 # show regression result
4 summary(Poly_fit)
5 # creat points to draw fitting line

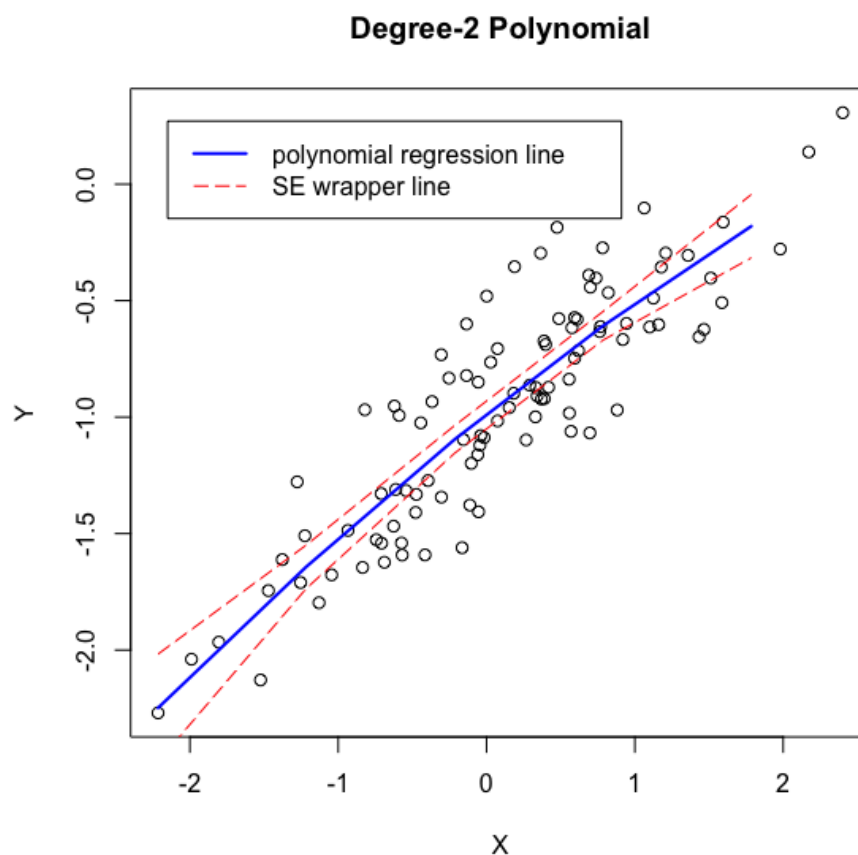
```

```

6 | xlims = range(x)
7 | x.grid = seq(from=xlims[1], to=xlims[2])
8 | preds = predict(Poly_fit, newdata=list(x=x.grid), se=TRUE)
9 | # use standard error to creat wrapper line
10 | se.bands = cbind(preds$fit+2*preds$se.fit, preds$fit-2*preds$se.fit)
11 | # draw
12 | plot(x, y, xlab="X", ylab="Y", main='Degree-2 Polynomial', col="black"
13 |      )
14 | lines(x.grid, preds$fit, lwd=2, col="blue", lty=1)
15 | matlines(x.grid, se.bands, lwd=1, col="red", lty=5)
16 | legend('topleft', inset=0.05, c('polynomial regression line', 'SE_
17 |   wrapper line'), lwd=c(2, 1), lty=c(1, 5), col=c('blue', 'red'), bty
18 |   = "o")

```

Result:



```

1 | Call:
2 | lm(formula = y ~ poly(x, 2))
3 |
4 | Residuals:
5 |      Min       1Q   Median       3Q      Max

```

```

6 -0.4913 -0.1563 -0.0322 0.1451 0.5675
7
8 Coefficients:
9             Estimate Std. Error t value Pr(>|t|)
10 (Intercept) -0.95501    0.02395  -39.874  <2e-16 ***
11 poly(x, 2)1  4.46612    0.23951   18.647  <2e-16 ***
12 poly(x, 2)2 -0.33602    0.23951   -1.403    0.164
13 ———
14 Signif. codes:  0    ***    0.001    **    0.01    *    0.05    .
15                 0.1      1
16 Residual standard error: 0.2395 on 97 degrees of freedom
17 Multiple R-squared:  0.7828,    Adjusted R-squared:  0.7784
18 F-statistic: 174.8 on 2 and 97 DF,  p-value: < 2.2e-16

```

Ans:

There is limited evidence that the quadratic term improves the model fit.

The $\Pr(> |t|)$ of the quadratic term equals 0.164. So we still can say there's no relation between Y and X^2 with the probability equalling 0.164, which is not a small enough probability.

Also, from the figure above, the polynomial regression line is very close to the least squares line in linear model.

Therefore, the quadratic term hardly improves the model fit.

(h)**Code:**

```

1 # a
2 set.seed(1) # ensure consistent results
3 x=rnorm(100, mean=0, sd=1) # feature X
4
5 # b
6 eps=rnorm(100, mean=0, sd=0.1)
7
8 # c
9 y = -1 + 0.5*x + eps
10 length(y)
11
12 # d scatterplot
13 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
14
15 # e least squares linear model
16 lm.fit = lm(y~x)
17 summary(lm.fit)
18 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
19 abline(lm.fit)
20
21 # f


```

```

22 par(col='black')
23 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
24 abline(lm.fit, col='blue', lty=5) # least squares line
25 abline(a=-1, b=0.5, col='red', lty=1) # population regression line
26 legend('topleft', inset=0.05, c('least_squares_line', 'population_
    regression_line'), lty=c(5, 1), col=c('blue', 'red'), bty = "o")

```

Result:



```

1 Call:
2 lm(formula = y ~ x)
3
4 Residuals:
5      Min       1Q   Median       3Q      Max
6 -0.18768 -0.06138 -0.01395  0.05394  0.23462
7
8 Coefficients:
9              Estimate Std. Error t value Pr(>|t|)
10 (Intercept) -1.003769   0.009699  -103.5  <2e-16 ***
11 x             0.499894   0.010773   46.4  <2e-16 ***
12 ———

```

```

13 Signif. codes:  0      ***      0.001      **      0.01      *      0.05      .
    0.1          1
14
15 Residual standard error: 0.09628 on 98 degrees of freedom
16 Multiple R-squared:  0.9565,    Adjusted R-squared:  0.956
17 F-statistic: 2153 on 1 and 98 DF,  p-value: < 2.2e-16

```

Ans:

With less noise in the data, the observation points are more near to the population regression line. The estimated value of β_0 and β_1 are more close to their true value. $R^2 = 0.9565$ means the linear model gives a better fit.

(i)**Code:**

```

1 # a
2 set.seed(1) # ensure consistent results
3 x=rnorm(100, mean=0, sd=1) # feature X
4
5 # b
6 eps=rnorm(100, mean=0, sd=1)
7
8 # c
9 y = -1 + 0.5*x + eps
10 length(y)
11
12 # d scatterplot
13 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
14
15 # e least squares linear model
16 lm.fit = lm(y~x)
17 summary(lm.fit)
18 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
19 abline(lm.fit)
20
21 # f
22 par(col='black')
23 plot(x, y, xlab="X", ylab="Y", main="Relationship between X and Y")
24 abline(lm.fit, col='blue', lty=5) # least squares line
25 abline(a=-1, b=0.5, col='red', lty=1) # population regression line
26 legend('topleft', inset=0.05, c('least_squares_line', 'population_
    regression_line'), lty=c(5, 1), col=c('blue', 'red'), bty = "o")

```

Result:

```


1 Call:
2 lm(formula = y ~ x)

```

```

3
4 Residuals:
5      Min       1Q   Median       3Q      Max
6 -1.8768 -0.6138 -0.1395  0.5394  2.3462
7
8 Coefficients:
9             Estimate Std. Error t value Pr(>|t|)
10 (Intercept) -1.03769    0.09699  -10.699  < 2e-16 ***
11 x            0.49894    0.10773   4.632  1.12e-05 ***
12 ---
13 Signif. codes:  0      ***    0.001    **    0.01    *    0.05    .
14                  0.1      1
15 Residual standard error: 0.9628 on 98 degrees of freedom
16 Multiple R-squared:  0.1796,    Adjusted R-squared:  0.1712
17 F-statistic: 21.45 on 1 and 98 DF,  p-value: 1.117e-05

```



Ans:

With more noise in the data, the observation points are more far to the population regression line. The estimated value of β_0 and β_1 are more different to their true value. $R^2 = 0.1796$ means the

linear model gives a worse fit.

(j)

Code:

```
1 confint(lm.fit)
```

Result:

$\epsilon \sim N(0, 0.25)$

```
1           2.5 %      97.5 %
2 (Intercept) -1.0575402 -0.9613061
3 x           0.4462897  0.5531801
```

$\epsilon \sim N(0, 0.1)$

```
1           2.5 %      97.5 %
2 (Intercept) -1.0230161 -0.9845224
3 x           0.4785159  0.5212720
```

$\epsilon \sim N(0, 1)$

```
1           2.5 %      97.5 %
2 (Intercept) -1.2301607 -0.8452245
3 x           0.2851588  0.7127204
```

Ans:

With less noise, the confidence interval for β_0 and β_1 are shorter, meaning the estimation are more accurate to the true value.

With more noise, the confidence interval for β_0 and β_1 are longer, meaning the estimation are less accurate to the true value.

P4

Code:

```

1 Advertising = read.csv("/Users/yangchenye/Downloads/Advertising.csv",
  header=T, na.strings="?")
2 dim(Advertising)
3 Advertising=na.omit(Advertising) # remove incomplete cases
4 dim(Advertising)
5 names(Advertising)
6 attach(Advertising)
7
8 # TV
9 TV_fit = lm(sales~TV) # least squares model
10 TV_CI = confint(TV_fit, level = 0.92) # 92% confidence intervals
11
12 plot(TV, sales, xlab="TV_advertising", ylab="Sales", main="
  Relationship_between_Sales_and_TV_advertising") # scatterplot
13 abline(TV_fit, col='blue', lty=1) # least squares line
14 abline(a=TV_CI[1,1], b=TV_CI[2,1], col='red', lty=5) # 92% confidence
  intervals line
15 abline(a=TV_CI[1,2], b=TV_CI[2,2], col='red', lty=5) # 92% confidence
  intervals line
16 legend('topleft', inset=0.05, c('least_squares_line', '92%_confidence_
  intervals_line'), lty=c(1, 5), col=c('blue', 'red'), bty = "o")
17
18 # Radio
19 Radio_fit = lm(sales~radio) # least squares model
20 Radio_CI = confint(Radio_fit, level = 0.92) # 92% confidence intervals
21
22 plot(radio, sales, xlab="Radio_advertising", ylab="Sales", main="
  Relationship_between_Sales_and_Radio_advertising") # scatterplot
23 abline(Radio_fit, col='blue', lty=1) # least squares line
24 abline(a=Radio_CI[1,1], b=Radio_CI[2,1], col='red', lty=5) # 92%
  confidence intervals line
25 abline(a=Radio_CI[1,2], b=Radio_CI[2,2], col='red', lty=5) # 92%
  confidence intervals line
26 legend('topleft', inset=0.05, c('least_squares_line', '92%_confidence_
  intervals_line'), lty=c(1, 5), col=c('blue', 'red'), bty = "o")
27
28 # Newspaper
29 Newspaper_fit = lm(sales~newspaper) # least squares model
30 Newspaper_CI = confint(Newspaper_fit, level = 0.92) # 92% confidence
  intervals
31
32 plot(newspaper, sales, xlab="Newspaper_advertising", ylab="Sales",
  main="Relationship_between_Sales_and_Newspaper_advertising") #

```

```

scatterplot
33 abline(Newspaper_fit , col='blue' , lty=1) # least squares line
34 abline(a=Newspaper_CI[1,1] , b=Newspaper_CI[2,1] , col='red' , lty=5) #
  92% confidence intervals line
35 abline(a=Newspaper_CI[1,2] , b=Newspaper_CI[2,2] , col='red' , lty=5) #
  92% confidence intervals line
36 legend('bottomright' , inset=0.05 , c('least_squares_line' , '92%_
  confidence_intervals_line') , lty=c(1 , 5) , col=c('blue' , 'red') , bty
    = "o")

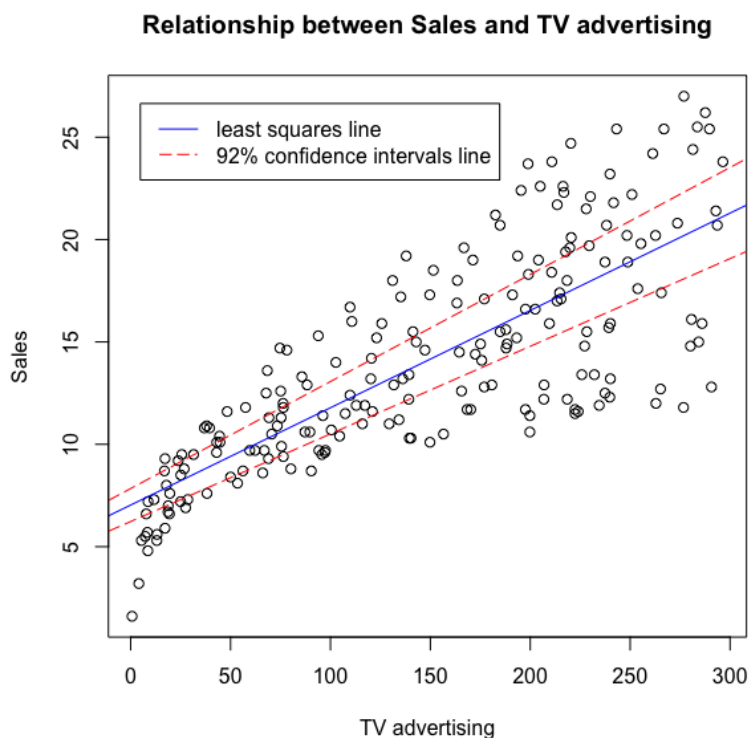
```


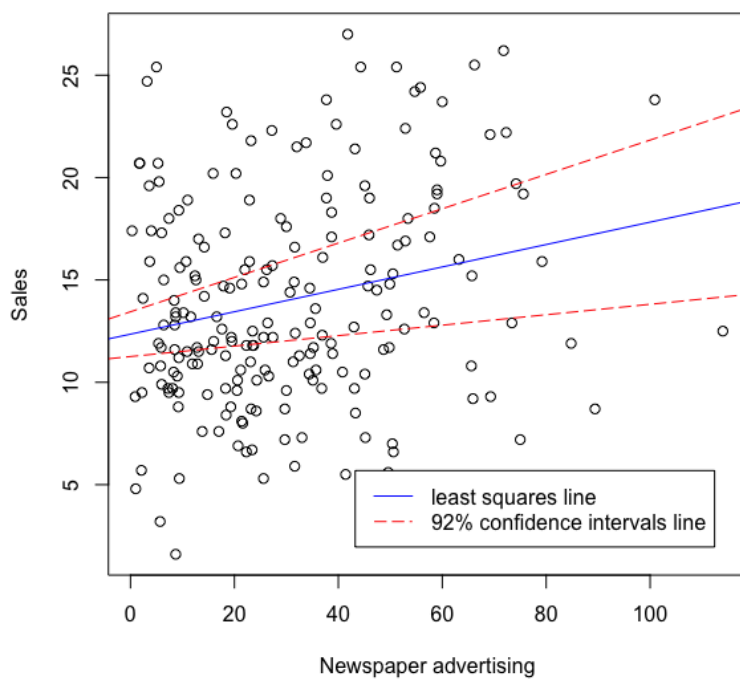
Result:

```

1 > TV_CI
2           4 %           96 %
3 (Intercept) 6.22691926 7.83826784
4 TV          0.04280193 0.05227135
5 > Radio_CI
6           4 %           96 %
7 (Intercept) 8.3210922 10.3021840
8 radio       0.1665776 0.2384139
9 > Newspaper_CI
10           4 %           96 %
11 (Intercept) 11.25788302 13.44493112
12 newspaper   0.02552451 0.08386169

```



Relationship between Sales and Radio advertising**Relationship between Sales and Newspaper advertising**

P5

(a)

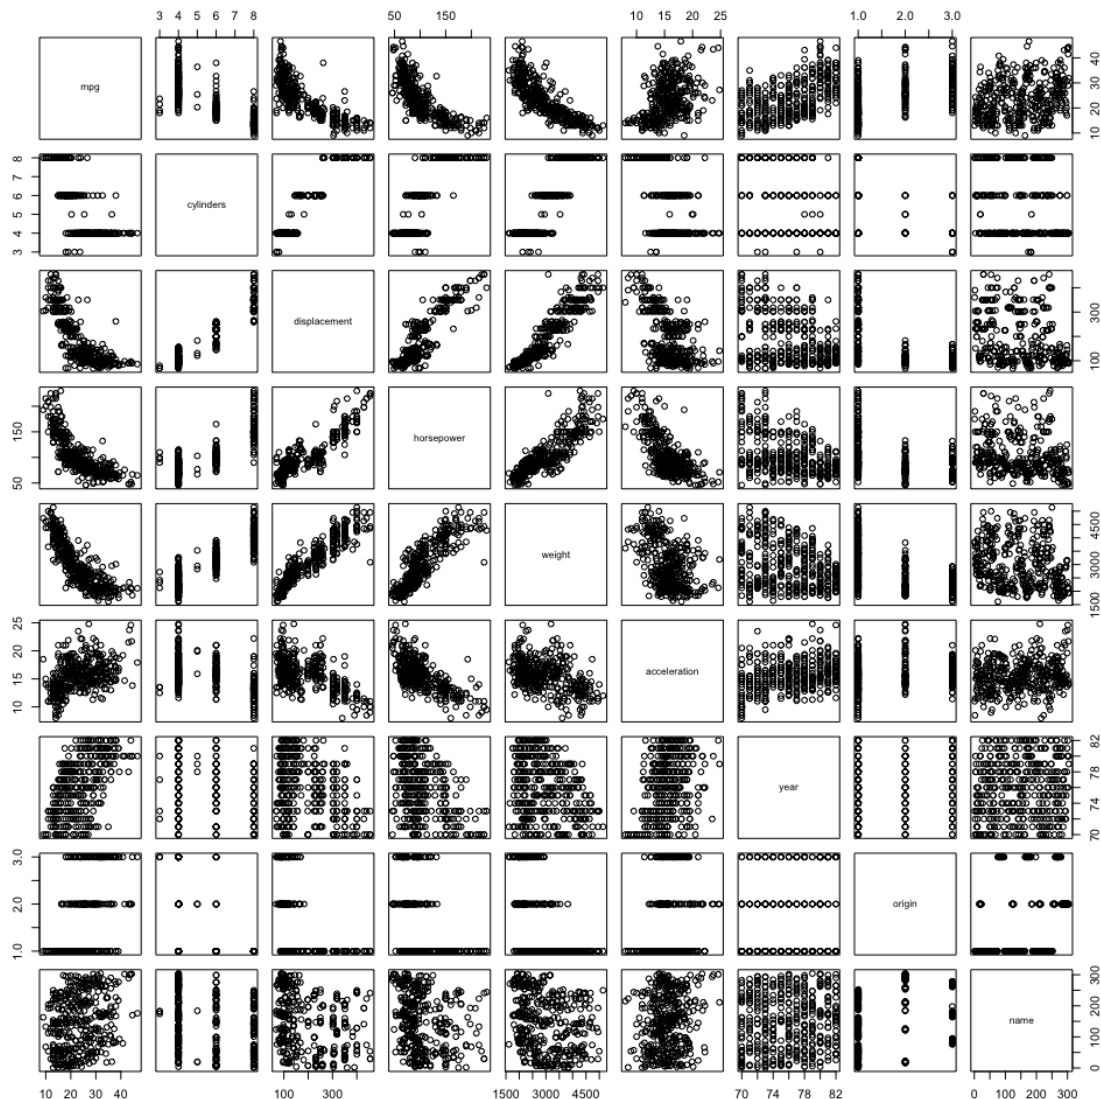
Code:

```

1 Auto = read.csv("/Users/yangchenye/Downloads/Auto.csv", header=T, na.
    strings="?")
2 dim(Auto)
3 Auto=na.omit(Auto) # remove incomplete cases
4 dim(Auto)
5 names(Auto)
6 attach(Auto)
7 pairs(Auto)

```

Result:



(b)

Code:

```
1 CRM = cor(Auto[,1:8])
```

Result:

```
1      mpg  cylinders displacement horsepower    weight acceleration    year    origin
2 mpg      1.0000000 -0.7776175  -0.8051269 -0.7784268 -0.8322442   0.4233285  0.5805410  0.5652088
3 cylinders -0.7776175  1.0000000   0.9508233  0.8429834  0.8975273  -0.5046834 -0.3456474 -0.5689316
4 displacement -0.8051269  0.9508233   1.0000000  0.8972570  0.9329944  -0.5438005 -0.3698552 -0.6145351
5 horsepower -0.7784268  0.8429834   0.8972570  1.0000000  0.8645377  -0.6891955 -0.4163615 -0.4551715
6 weight      -0.8322442  0.8975273   0.9329944  0.8645377  1.0000000  -0.4168392 -0.3091199 -0.5850054
7 acceleration 0.4233285 -0.5046834  -0.5438005 -0.6891955 -0.4168392  1.0000000  0.2903161  0.2127458
8 year        0.5805410 -0.3456474  -0.3698552 -0.4163615 -0.3091199  0.2903161  1.0000000  0.1815277
9 origin      0.5652088 -0.5689316  -0.6145351 -0.4551715 -0.5850054  0.2127458  0.1815277  1.0000000
```

(c)

Code:

```
1 mlm_fit = lm(mpg~cylinders + displacement + horsepower + weight +
2             acceleration + year + origin)
3 summary(mlm_fit)
```

Result:

```
1 Call:
2 lm(formula = mpg ~ cylinders + displacement + horsepower + weight +
3     acceleration + year + origin)
4
5 Residuals:
6      Min       1Q   Median       3Q      Max
7 -9.5903 -2.1565 -0.1169  1.8690 13.0604
8
9 Coefficients:
10             Estimate Std. Error t value Pr(>|t|)
11 (Intercept) -17.218435   4.644294  -3.707  0.00024 ***
12 cylinders    -0.493376   0.323282  -1.526  0.12780
13 displacement  0.019896   0.007515   2.647  0.00844 **
14 horsepower   -0.016951   0.013787  -1.230  0.21963
15 weight       -0.006474   0.000652  -9.929 < 2e-16 ***
16 acceleration  0.080576   0.098845   0.815  0.41548
17 year         0.750773   0.050973  14.729 < 2e-16 ***
18 origin       1.426141   0.278136   5.127 4.67e-07 ***
19
20 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
21                 0.1 '1'
22 Residual standard error: 3.328 on 384 degrees of freedom
23 Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
24 F-statistic: 252.4 on 7 and 384 DF, p-value: < 2.2e-16
```

Ans:

i. Is there a relationship between the predictors and the response?

With the p-value $< 2.2e - 16$, we should reject H_0 , namely, admit that there exist a relationship between predictors and response. And $R^2 = 0.8215$ shows that this linear model fits the data well.

ii. Which predictors appear to have a statistically significant relationship to the response?

The $\Pr(> |t|)$ of cylinders, horsepower and acceleration are not small enough to say they have significant relationship to the response, while the $\Pr(> |t|)$ of other predictors are small enough. That is to say, displacement, weight, year and origin have a statistically significant relationship to the response mpg.

iii. What does the coefficient for the year variable suggest?

The estimated coefficient for the year variable is 0.750773, suggesting that with one unit increase in year, the mpg will increase about 0.75 unit. The year has a positive relationship to mpg.

(d)

Code:

```

1 mlm_fit_log = lm(log10(mpg)~log10(cylinders) + log10(displacement) +
  log10(horsepower) + log10(weight) + log10(acceleration) + log10(
  year) + log10(origin))
2 summary(mlm_fit_log)
3
4 mlm_fit_sqrt = lm(sqrt(mpg)~sqrt(cylinders) + sqrt(displacement) +sqrt
  (horsepower) + sqrt(weight) + sqrt(acceleration) + sqrt(year) +
  sqrt(origin))
5 summary(mlm_fit_sqrt)
6
7 mlm_fit_power = lm(mpg^2~cylinders^2 + displacement^2 + horsepower^2 +
  weight^2 + acceleration^2 + year^2 + origin^2)
8 summary(mlm_fit_power)

```

Result with the transformation of $\log(x)$:

```

1 Call:
2 lm(formula = log10(mpg) ~ log10(cylinders) + log10(displacement) +
3   log10(horsepower) + log10(weight) + log10(acceleration) +
4   log10(year) + log10(origin))
5
6 Residuals:
7      Min       1Q   Median       3Q      Max
8 -0.179356 -0.030825  0.000238  0.026708  0.171685
9
10 Coefficients:
11             Estimate Std. Error t value Pr(>|t|)
12 (Intercept)   -0.067485   0.281523  -0.240  0.81068
13 log10(cylinders) -0.082815   0.061429  -1.348  0.17841
14 log10(displacement)  0.006625   0.056970   0.116  0.90748

```

```

15 log10(horsepower)    -0.294389    0.057652    -5.106  5.18e-07 ***
16 log10(weight)       -0.569666    0.082397    -6.914  1.98e-11 ***
17 log10(acceleration) -0.179239    0.059536    -3.011  0.00278 **
18 log10(year)         2.243989    0.131661    17.044  < 2e-16 ***
19 log10(origin)       0.044848    0.018821     2.383  0.01767 *
20
21 Signif. codes:  0      ***      0.001      **      0.01      *      0.05      .
    0.1              1
22
23 Residual standard error: 0.04935 on 384 degrees of freedom
24 Multiple R-squared:  0.8903,    Adjusted R-squared:  0.8883
25 F-statistic: 445.3 on 7 and 384 DF,  p-value: < 2.2e-16

```

Result with the transformation of \sqrt{x} :

```

1 Call:
2 lm(formula = sqrt(mpg) ~ sqrt(cylinders) + sqrt(displacement) +
3     sqrt(horsepower) + sqrt(weight) + sqrt(acceleration) + sqrt(year)
4     +
5     sqrt(origin))
6 Residuals:
7      Min       1Q   Median       3Q      Max
8 -0.98667 -0.17280 -0.00315  0.16145  1.02245
9
10 Coefficients:
11             Estimate Std. Error t value Pr(>|t|)
12 (Intercept)   -1.949286   0.847481  -2.300  0.021979 *
13 sqrt(cylinders) -0.108552   0.141968  -0.765  0.444964
14 sqrt(displacement) 0.019707   0.021182   0.930  0.352752
15 sqrt(horsepower) -0.090896   0.028428  -3.197  0.001502 **
16 sqrt(weight)    -0.061414   0.007292  -8.422  7.48e-16 ***
17 sqrt(acceleration) -0.107258   0.077048  -1.392  0.164699
18 sqrt(year)      1.266015   0.079308  15.963  < 2e-16 ***
19 sqrt(origin)    0.272324   0.070883   3.842  0.000143 ***
20
21 Signif. codes:  0      ***      0.001      **      0.01      *      0.05      .
    0.1              1
22
23 Residual standard error: 0.2964 on 384 degrees of freedom
24 Multiple R-squared:  0.8662,    Adjusted R-squared:  0.8638
25 F-statistic: 355.1 on 7 and 384 DF,  p-value: < 2.2e-16

```

Result with the transformation of x^2 :

```

1 Call:
2 lm(formula = mpg^2 ~ cylinders^2 + displacement^2 + horsepower^2 +
3     weight^2 + acceleration^2 + year^2 + origin^2)

```



```

4
5 Residuals :
6      Min       1Q   Median       3Q      Max
7 -483.45 -141.87  -19.62   103.58 1042.84
8
9 Coefficients :
10              Estimate Std. Error t value Pr(>|t|)
11 (Intercept) -1.878e+03  2.928e+02  -6.412 4.22e-10 ***
12 cylinders   -1.436e+01  2.038e+01  -0.704  0.48157
13 displacement  1.328e+00  4.738e-01   2.802  0.00534 **
14 horsepower   -3.587e-01  8.693e-01  -0.413  0.68009
15 weight       -3.522e-01  4.111e-02  -8.567 2.62e-16 ***
16 acceleration  9.278e+00  6.232e+00   1.489  0.13740
17 year         4.081e+01  3.214e+00  12.698 < 2e-16 ***
18 origin       9.509e+01  1.754e+01   5.422 1.04e-07 ***
19 ———
20 Signif. codes:  0      ***      0.001      **      0.01      *      0.05      .
21                  0.1          1
22 Residual standard error: 209.8 on 384 degrees of freedom
23 Multiple R-squared:  0.7292,    Adjusted R-squared:  0.7243
24 F-statistic: 147.8 on 7 and 384 DF,  p-value: < 2.2e-16

```

Ans:

Different transformations of variables will lead to different estimate value of parameters in a same linear model. For example, the intercepts corresponding with different transformation are 17.218435, 0.067485, 1.949286 and 1.878e+03.

Also, transformations will change the relationship between variables. For example, "displacement" has relationship with "mpg" in x and x^2 transformation situation, while little relationship in $\log(x)$ and \sqrt{x} situation.

Moreover, transformation on variables will also change the property of the linear model, which can be seen in the difference of R^2 , RSE and F - statistic.

P6

Ans:

From the given constrain, we can calculate the following values:

$$\begin{aligned}\bar{x} &= \frac{1}{n} \sum_{i=1}^n x_i = \frac{1}{20} \sum_{i=1}^{20} x_i = \frac{1}{20} \times 8.552 = 0.4276 \\ \bar{y} &= \frac{1}{n} \sum_{i=1}^n y_i = \frac{1}{20} \sum_{i=1}^{20} y_i = \frac{1}{20} \times 398.2 = 19.91\end{aligned}\tag{18}$$

$$\begin{aligned}\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) &= \sum_{i=1}^{20} (x_i y_i - \bar{x} y_i - x_i \bar{y} + \bar{x} \bar{y}) \\ &= \sum_{i=1}^{20} (x_i y_i) - \bar{x} \sum_{i=1}^{20} y_i - \bar{y} \sum_{i=1}^{20} x_i + 20 \bar{x} \bar{y} \\ &= 216.6 - 0.4276 \times 398.2 - 19.91 \times 8.552 + 20 \times 0.4276 \times 19.91 \\ &= 46.33\end{aligned}$$

$$\begin{aligned}\sum_{i=1}^n (x_i - \bar{x})^2 &= \sum_{i=1}^{20} (x_i^2 - 2\bar{x}x_i + \bar{x}^2) \\ &= \sum_{i=1}^{20} x_i^2 - 2\bar{x} \sum_{i=1}^{20} x_i + 20\bar{x}^2 \\ &= 5.196 - 2 \times 0.4276 \times 8.552 + 20 \times 0.4276^2 \\ &= 1.539\end{aligned}\tag{19}$$

$$\begin{aligned}\sum_{i=1}^n (y_i - \bar{y})^2 &= \sum_{i=1}^{20} (y_i^2 - 2\bar{y}y_i + \bar{y}^2) \\ &= \sum_{i=1}^{20} y_i^2 - 2\bar{y} \sum_{i=1}^{20} y_i + 20\bar{y}^2 \\ &= 9356 - 2 \times 19.91 \times 398.2 + 20 \times 19.91^2 \\ &= 1428\end{aligned}$$

Thus:

$$\begin{aligned}\hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{46.33}{1.539} = 30.10 \\ \hat{\beta}_0 &= \bar{y} - \hat{\beta}_1 \bar{x} = 19.91 - 30.10 \times 0.4276 = 7.039 \\ \text{RSS} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2 = \sum_{i=1}^n (y_i - \bar{y} + \hat{\beta}_1 \bar{x} - \hat{\beta}_1 x_i)^2 \\ &= \sum_{i=1}^n (y_i - \bar{y})^2 - \hat{\beta}_1 \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x}) = 1428 - 30.10 \times 46.33 = 33.47\end{aligned}\tag{20}$$

$$\begin{aligned} R^2 &= 1 - \frac{\text{RSS}}{\sum_{i=1}^n (y_i - \bar{y})^2} = 1 - \frac{33.47}{1428} = 0.9766 \\ \hat{\sigma}^2 &= \frac{\text{RSS}}{n - 2} = \frac{33.47}{20 - 2} = 1.859 \end{aligned} \tag{21}$$

When $x = 0.5$:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x = 7.039 + 30.10 \times 0.5 = 22.09$$

P7

Code:

```
1 fcdf(1.89, 6, 38, 'upper') % matlab
```

Ans:

The null hypothesis test is performed by computing the F-statistic:

$$F = \frac{(\text{TSS} - \text{RSS})/p}{\text{RSS}/(n - p - 1)} = \frac{(11.62 - 8.95)/6}{8.95/(45 - 6 - 1)} = 1.89 \quad (22)$$
$$\Pr(F_{6,38} > 1.89) = 0.11$$

The p -value for the null hypothesis is 0.11

E1**Ans:**

The definition of gamma function is :

$$\Gamma(x) = \int_0^{+\infty} t^{x-1} e^{-t} dt$$

If a random variable $X \sim \Gamma(\alpha, \theta)$, then the probability density function of X has the following format:

$$f_X(x) = \begin{cases} \frac{1}{\theta^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\theta}, & x > 0 \\ 0 & , \text{else} \end{cases} \quad (23)$$

Let the density of X be $f_X(x)$, and let the density of $Y = X^2$ be $f_Y(y)$, we have:

$$f_Y(y) = \begin{cases} \frac{1}{2\sqrt{y}} [f_X(\sqrt{y}) + f_X(-\sqrt{y})], & y > 0 \\ 0 & , \text{else} \end{cases} \quad (24)$$

Therefore, as for random variable $X \sim N(0, 1)$ with density $\varphi(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}$ ($-\infty < x < \infty$), let $Y = X^2$, the density of Y is:

$$\begin{aligned} f_Y(y) &= \begin{cases} \frac{1}{2\sqrt{y}} \left[\frac{1}{\sqrt{2\pi}} e^{-y/2} + \frac{1}{\sqrt{2\pi}} e^{-y/2} \right], & y > 0 \\ 0 & , \text{else} \end{cases} \\ &= \begin{cases} \frac{1}{\sqrt{2\pi}} \frac{1}{\sqrt{y}} e^{-y/2}, & y > 0 \\ 0 & , \text{else} \end{cases} \end{aligned} \quad (25)$$

Thus, $Y = X^2 \sim \Gamma(\frac{1}{2}, 2)$.

Because X_i , $i = 1, \dots, n$ are independent, and as a result X_i^2 , $i = 1, \dots, n$ are also independent, Chi-squared distribution $\chi_n^2 = X_1^2 + X_2^2 + \dots + X_n^2$ has additivity. Namely:

$$\chi_n^2 = \sum_{i=1}^n X_i^2 \sim \Gamma\left(\frac{n}{2}, 2\right)$$

Therefore, the density of χ_n^2 is given by:

$$g_n(x) = \begin{cases} \frac{1}{2^{n/2} \Gamma(n/2)} x^{n/2-1} e^{-x/2}, & x > 0 \\ 0 & , x \leq 0 \end{cases} \quad (26)$$

E2**Ans:**

(The first half of the proof is the same as P1(d), I just copy and paste.)

Let $Z_i = \frac{X_i - \mu}{\sigma}$ $i = 1, 2, \dots, n$, thus Z_1, Z_2, \dots, Z_n are independent and $Z_i \sim N(0, 1)$ $i = 1, 2, \dots, n$.
Also:

$$\begin{aligned}
 \bar{Z} &= \frac{1}{n} \sum_{i=1}^n Z_i = \frac{1}{n} \sum_{i=1}^n \frac{X_i - \mu}{\sigma} = \frac{\bar{X} - \mu}{\sigma} \\
 \frac{(n-1)S^2}{\sigma^2} &= \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{\sigma^2} = \sum_{i=1}^n \left[\frac{(X_i - \mu) - (\bar{X} - \mu)}{\sigma} \right]^2 \\
 &= \sum_{i=1}^n (Z_i - \bar{Z})^2 = \sum_{i=1}^n (Z_i^2 - 2Z_i\bar{Z} + \bar{Z}^2) \\
 &= \sum_{i=1}^n Z_i^2 - 2\bar{Z} \sum_{i=1}^n Z_i + n\bar{Z}^2 = \sum_{i=1}^n Z_i^2 - n\bar{Z}^2
 \end{aligned} \tag{27}$$

Choose a matrix $\mathbf{A} = (a_{ij})$, in which all the entries in the first row are $1/\sqrt{n}$, and each row vector of the matrix should be orthogonal. Do:

$$\mathbf{Y} = \mathbf{AZ}$$

In which

$$\mathbf{Y} = (Y_1, Y_2, \dots, Y_n)^T$$

$$\mathbf{Z} = (Z_1, Z_2, \dots, Z_n)^T$$

$$\mathbf{A} = (\mathbf{a}_1^T, \mathbf{a}_2^T, \dots, \mathbf{a}_n^T)^T$$

Because Y_i is the linear combination of Z_i and $Z_i \sim N(0, 1)$ $i = 1, 2, \dots, n$, therefore Y_i $i = 1, 2, \dots, n$ have normal/Gaussian distribution and:

$$\begin{aligned}
 \mathbb{E}(Y_i) &= \mathbb{E}\left(\sum_{j=1}^n a_{ij} Z_j\right) = \sum_{j=1}^n a_{ij} \mathbb{E}(Z_j) = 0 \\
 \text{Var}(Y_i) &= \text{Var}\left(\sum_{j=1}^n a_{ij} Z_j\right) = \sum_{j=1}^n a_{ij}^2 = \langle \mathbf{a}_i, \mathbf{a}_i \rangle = 1
 \end{aligned} \tag{28}$$

Because $\text{Cov}(Z_k, Z_h)$ equals 1 when $k = h$ and equals 0 when $k \neq h$:

$$\begin{aligned}
 \text{Cov}(Y_i, Y_j) &= \text{Cov}\left(\sum_{k=1}^n a_{ik} Z_k, \sum_{h=1}^n a_{jh} Z_h\right) \\
 &= \sum_{k=1}^n \sum_{h=1}^n a_{ik} a_{jh} \text{Cov}(Z_k, Z_h) \\
 &= \sum_{k=1}^n a_{ik} a_{jk} \\
 &= \langle \mathbf{a}_i, \mathbf{a}_j \rangle \\
 &= \begin{cases} 1 & , i = j \\ 0 & , i \neq j \end{cases} \\
 &= \text{Cov}(Z_i, Z_j)
 \end{aligned} \tag{29}$$

Therefore, Y_i $i = 1, 2, \dots, n$ has no correlation with each other. Considering $Y_i \sim N(0, 1)$, thus Y_1, Y_2, \dots, Y_n are independent. Also:

$$\begin{aligned}
 Y_1 &= \sum_{j=1}^n a_{1j} Z_j = \sum_{j=1}^n \frac{1}{\sqrt{n}} Z_j = \sqrt{n} \bar{Z} \\
 \sum_{i=1}^n Y_i^2 &= \mathbf{Y}^T \mathbf{Y} = (\mathbf{A}\mathbf{Z})^T (\mathbf{A}\mathbf{Z}) = \mathbf{Z}^T (\mathbf{A}\mathbf{A}^T)^T \mathbf{Z} = \mathbf{Z}^T \mathbf{I} \mathbf{Z} = \sum_{i=1}^n Z_i^2
 \end{aligned} \tag{30}$$

Therefore:

$$\frac{(n-1)S^2}{\sigma^2} = \sum_{i=1}^n Z_i^2 - n\bar{Z}^2 = \sum_{i=1}^n (Y_i^2) - Y_1^2 = \sum_{i=2}^n (Y_i^2) \tag{31}$$

Because Y_2, Y_3, \dots, Y_n are independent and $Y_i \sim N(0, 1)$ $i = 2, 3, \dots, n$, thus $\sum_{i=2}^n (Y_i^2) \sim \chi_{n-1}^2$. Thus:

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2$$

Namely:

$$\frac{(n-1)S^2}{\sigma^2} \stackrel{d}{=} \chi_{n-1}^2$$

E3**Ans:**

E4**Ans:**