

Statistical Learning for Biological and Information Systems

Problem Set #3

Chenye Yang cy2540@columbia.edu

October 22, 2019

P1

(a)

Ans:

Assume A = "receive an A", \bar{A} = "not receive an A". Thus:

$$\begin{aligned}\log \frac{p_A(X)}{p_{\bar{A}}(X)} &= \beta_0 + \beta_1 X_1 + \beta_2 X_2 \\ &= -6 + 0.05X_1 + X_2\end{aligned}$$

For $X_1 = 40$, $X_2 = 3.5$,

$$\log \frac{p_A(X)}{1 - p_{\bar{A}}(X)} = -6 + 0.05 \times 40 + 3.5 = -0.5$$

\Rightarrow

$$p_A(X) = 0.38$$

Therefore, the probability that a student who studies for 40 hours and has an undergrad GPA of 3.5 gets an A in the class is 0.38

(b)

Ans:

When $p_A(X) = 0.5$, $\log \frac{p_A(X)}{p_{\bar{A}}(X)} = 0$

$$\begin{cases} -6 + 0.05X_1 + X_2 = 0 \\ X_2 = 3.5 \end{cases}$$

\Rightarrow

$$X_1 = 50$$

Therefore, the student in part (a) need to study 50 hours to have a 50% chance of getting an A in the class.

P2**Ans:**

Given:

$$Y = \text{Yes} \quad \bar{X} = 10 \quad \hat{\sigma}^2 = 36 \quad \pi_{\text{Yes}} = 0.8 \quad X \sim N(10, 36)$$

$$Y = \text{No} \quad \bar{X} = 0 \quad \hat{\sigma}^2 = 36 \quad \pi_{\text{No}} = 0.2 \quad X \sim N(0, 36)$$

Thus,

$$\begin{aligned} P(Y = \text{Yes} | X = 4) &= \frac{\pi_{\text{Yes}} f_{\text{Yes}}(X)}{\pi_{\text{Yes}} f_{\text{Yes}}(X) + \pi_{\text{No}} f_{\text{No}}(X)} \\ &= \frac{0.8 \times \frac{1}{6\sqrt{2\pi}} e^{-\frac{(4-10)^2}{2 \times 36}}}{0.8 \times \frac{1}{6\sqrt{2\pi}} e^{-\frac{(4-10)^2}{2 \times 36}} + 0.2 \times \frac{1}{6\sqrt{2\pi}} e^{-\frac{(4-0)^2}{2 \times 36}}} \\ &= \frac{4 \times e^{-1/2}}{4 \times e^{-1/2} + e^{-2/9}} \\ &= 0.752 \end{aligned}$$

Therefore, the probability that a company will issue a dividend this year given that its percentage profit was $X = 4$ last year is 0.752

P3

Ans:

Assume $\mathbf{X}_1 = [1, 0.0]^T$, $\mathbf{X}_2 = [1, 0.2]^T$, $\mathbf{X}_3 = [1, 0.4]^T$, $\mathbf{X}_4 = [1, 0.6]^T$, $\mathbf{X}_5 = [1, 0.8]^T$, $\mathbf{X}_6 = [1, 1.0]^T$ and $y_1 = 0$, $y_2 = 0$, $y_3 = 0$, $y_4 = 1$, $y_5 = 0$, $y_6 = 1$ and $\boldsymbol{\beta} = [\beta_0, \beta_1]^T$.

The log-likelihood function, $\ell(\boldsymbol{\beta}) = \ell(\beta_0, \beta_1)$:

$$\begin{aligned}
 \ell(\boldsymbol{\beta}) &= \sum_{i=1}^6 [y_i \log p(\mathbf{X}_i; \boldsymbol{\beta}) + (1 - y_i) \log (1 - p(\mathbf{X}_i; \boldsymbol{\beta}))] \\
 &= \sum_{i=1}^6 \left[y_i \boldsymbol{\beta}^T \mathbf{X}_i - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_i}) \right] \\
 &= -\log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_1}) - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_2}) - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_3}) + \boldsymbol{\beta}^T \mathbf{X}_4 - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_4}) \\
 &\quad - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_5}) + \boldsymbol{\beta}^T \mathbf{X}_6 - \log(1 + e^{\boldsymbol{\beta}^T \mathbf{X}_6}) \\
 &= -\log(1 + e^{\beta_0}) - \log(1 + e^{\beta_0 + \beta_1 \times 0.2}) - \log(1 + e^{\beta_0 + \beta_1 \times 0.4}) + (\beta_0 + \beta_1 \times 0.6) \\
 &\quad - \log(1 + e^{\beta_0 + \beta_1 \times 0.6}) - \log(1 + e^{\beta_0 + \beta_1 \times 0.8}) + (\beta_0 + \beta_1 \times 1.0) - \log(1 + e^{\beta_0 + \beta_1 \times 1.0})
 \end{aligned} \tag{1}$$

To maximize the log-likelihood, we set its derivatives to zero:

$$\frac{\partial \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = \sum_{i=1}^6 \mathbf{X}_i [y_i - p(\mathbf{X}_i; \boldsymbol{\beta})] = 0 \tag{2}$$

In practice, we use the Newton–Raphson algorithm to solve the maximizing problem:

$$\begin{aligned}
 \frac{\partial \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} &= \sum_{i=1}^6 \mathbf{X}_i [y_i - p(\mathbf{X}_i; \boldsymbol{\beta})] \\
 \frac{\partial^2 \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} &= - \sum_{i=1}^6 \mathbf{X}_i \mathbf{X}_i^T p(\mathbf{X}_i; \boldsymbol{\beta}) (1 - p(\mathbf{X}_i; \boldsymbol{\beta})) \\
 \boldsymbol{\beta}^{\text{new}} &= \boldsymbol{\beta}^{\text{old}} - \left(\frac{\partial^2 \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} \right)^{-1} \frac{\partial \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}}
 \end{aligned} \tag{3}$$

Code:

```

1 beta_old = c(0,0)
2 beta_new = c(0,0)
3 # initiate X
4 X = matrix(data = rep(0,12), nrow = 2, ncol = 6)
5 for (i in 1:6){
6   X[,i] = c(1,(i-1)/5)
7 }
8 # initiate y
9 y = c(0, 0, 0, 1, 0, 1)
10 # initiate derivatives
11 l_b = c(0, 0)
12 for (i in 1:6){

```

```

13 l_b = l_b + X[,i]*(y[i]-exp(X[,i] %*% beta_old)/(1+exp(X[,i] %*%
    beta_old)))
14 }
15 # initiate second derivatives
16 l_bb = matrix(data = rep(0,4), nrow = 2, ncol = 2)
17 for (i in 1:6){
18   l_bb = l_bb - (X[,i] %*% t(X[,i]))*((exp(X[,i] %*% beta_old)/(1+exp(
    X[,i] %*% beta_old)))*(1-(exp(X[,i] %*% beta_old)/(1+exp(X[,i] %*%
    % beta_old))))) [1]
19 }
20 # iteration
21 for (i in 1:10){
22   beta_new = t(t(beta_old)) - solve(l_bb) %*% t(t(l_b))
23   beta_old = t(t(beta_new))
24 }
25 t(t(beta_new))

```

Result:

```

1 > t(t(beta_new))
2      [,1]
3 [1,] -23.80952
4 [2,]  34.28571

```

The estimated coefficients of this logistic regression problem is $\hat{\beta}_0 = -23.81$, $\hat{\beta}_1 = 34.29$.

P4**Ans:**

From given,

$$\begin{aligned}
\text{cov}[\mathbf{Y}] &= \text{cov}[\mathbf{A}\mathbf{X}] \\
&= \mathbf{A} \text{cov}[\mathbf{X}] \mathbf{A}^T \\
&= \mathbf{A} \mathbf{\Sigma} \mathbf{A}^T = \mathbf{I}
\end{aligned} \tag{4}$$

Because $\mathbf{\Sigma}$ is symmetric, assume $\mathbf{\Sigma}$ has non-degenerate eigenvalues λ_1, λ_2 and corresponding linearly independent eigenvectors v_1, v_2 . Define:

$$\begin{aligned}
\mathbf{Q} &= [\mathbf{v}_1 \quad \mathbf{v}_2] = \begin{bmatrix} v_1^1 & v_2^1 \\ v_1^2 & v_2^2 \end{bmatrix} \\
\mathbf{\Lambda} &= \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}
\end{aligned} \tag{5}$$

Thus we have,

$$\mathbf{\Sigma} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1} \tag{6}$$

Therefore,

$$\begin{aligned}
\mathbf{A} \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T \mathbf{A}^T &= \mathbf{I} \\
(\mathbf{A} \mathbf{Q}) \mathbf{\Lambda} (\mathbf{A} \mathbf{Q})^T &= \mathbf{I}
\end{aligned} \tag{7}$$

 \Rightarrow

$$\begin{aligned}
\mathbf{A} \mathbf{Q} &= (\mathbf{\Lambda}^{-1})^{\frac{1}{2}} \\
\mathbf{A} &= (\mathbf{\Lambda}^{-1})^{\frac{1}{2}} \mathbf{Q}^{-1}
\end{aligned} \tag{8}$$

Using Matlab program to calculate the eigenvalues and eigenvectors of $\mathbf{\Lambda}$:**Code:**

```

1 # Matlab
2 syms sigma_1 sigma_2 rho
3 A = [sigma_1^2, rho*sigma_1*sigma_2;
4      rho*sigma_1*sigma_2, sigma_2^2];
5 [eigenvector, eigenvalue] = eig(A)

```

Result:

$$\begin{aligned}
\mathbf{Q} &= [\mathbf{v}_1 \quad \mathbf{v}_2] \\
\mathbf{v}_1 &= \begin{bmatrix} (\sigma_1^2/2 - (4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2}/2 + \sigma_2^2/2)/(\rho\sigma_1\sigma_2) - \sigma_2/(\rho\sigma_1) \\ 1 \end{bmatrix} \\
\mathbf{v}_2 &= \begin{bmatrix} ((4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2}/2 + \sigma_1^2/2 + \sigma_2^2/2)/(\rho\sigma_1\sigma_2) - \sigma_2/(\rho\sigma_1) \\ 1 \end{bmatrix} \\
\mathbf{\Lambda} &= \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \\
\lambda_1 &= \sigma_1^2/2 - (4 * \rho^2 * \sigma_1^2 * \sigma_2^2 + \sigma_1^4 - 2 * \sigma_1^2 * \sigma_2^2 + \sigma_2^4)^{1/2}/2 + \sigma_2^2/2 \\
\lambda_2 &= (4 * \rho^2 * \sigma_1^2 * \sigma_2^2 + \sigma_1^4 - 2 * \sigma_1^2 * \sigma_2^2 + \sigma_2^4)^{1/2}/2 + \sigma_1^2/2 + \sigma_2^2/2
\end{aligned} \tag{9}$$

Therefore, use Matlab program to calculate \mathbf{A} :

Code:

```
1 A = (inv(eigenvalue))^(1/2)*inv(eigenvector)
```

Result:

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$$

$$a_{11} = -(2^{1/2}\rho\sigma_1\sigma_2(1/(\sigma_1^2 - (4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} + \sigma_2^2))^{1/2})/(4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2}$$

$$a_{12} = (2^{1/2}(1/(\sigma_1^2 - (4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} + \sigma_2^2))^{1/2}((4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} + \sigma_1^2 - \sigma_2^2))$$

$$/(2(4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2})$$

$$a_{21} = (2^{1/2}\rho\sigma_1\sigma_2(1/((4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} + \sigma_1^2 + \sigma_2^2))^{1/2})/(4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2}$$

$$a_{22} = (2^{1/2}(1/((4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} + \sigma_1^2 + \sigma_2^2))^{1/2}((4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2} - \sigma_1^2 + \sigma_2^2))$$

$$/(2(4\rho^2\sigma_1^2\sigma_2^2 + \sigma_1^4 - 2\sigma_1^2\sigma_2^2 + \sigma_2^4)^{1/2})$$

(10)

P5**Ans:**

Because the variance of each population is the same (σ^2), we have:

$$\begin{aligned}\hat{\sigma}_k^2 &= \frac{1}{n_k - 1} \sum_{i:y_i=k} (x_i - \hat{\mu}_k)^2 \\ \mathbb{E} [\hat{\sigma}_k^2] &= \sigma^2\end{aligned}\tag{11}$$

Thus,

$$\begin{aligned}\mathbb{E} [\hat{\sigma}^2] &= \mathbb{E} \left[\sum_{k=1}^K \alpha_k \hat{\sigma}_k^2 \right] = \sum_{k=1}^K \alpha_k \mathbb{E} [\hat{\sigma}_k^2] \\ &= \sigma^2 \sum_{k=1}^K \alpha_k = \sigma^2\end{aligned}\tag{12}$$

Because x_i obeys Gaussian distribution, thus

$$\frac{(n_k - 1) \hat{\sigma}_k^2}{\sigma^2} \sim \chi_{n_k - 1}^2$$

Therefore,

$$\begin{aligned}\text{var} \left[\frac{(n_k - 1) \hat{\sigma}_k^2}{\sigma^2} \right] &= 2(n_k - 1) \\ \frac{(n_k - 1)^2}{\sigma^4} \text{var} [\hat{\sigma}_k^2] &= 2(n_k - 1) \\ \text{var} [\hat{\sigma}_k^2] &= \frac{2\sigma^4}{n_k - 1}\end{aligned}\tag{13}$$

Considering,

$$\begin{aligned}\text{var} [\hat{\sigma}^2] &= \text{var} \left[\sum_{k=1}^K \alpha_k \hat{\sigma}_k^2 \right] \\ &= \sum_{k=1}^K \alpha_k^2 \text{var} [\hat{\sigma}_k^2] \\ &= \sum_{k=1}^K \alpha_k^2 \frac{2\sigma^4}{n_k - 1} \\ &= 2\sigma^4 \sum_{k=1}^K \frac{\alpha_k^2}{n_k - 1} \sum_{i=1}^K \alpha_i \\ &= 2\sigma^4 \sum_{k=1}^K \frac{\alpha_k^2}{n_k - 1} \sum_{i=1}^K \frac{n_i - 1}{n - K} \\ &= \frac{2\sigma^4}{n - K} \sum_{k=1}^K \frac{\alpha_k^2}{n_k - 1} \sum_{i=1}^K (n_i - 1)\end{aligned}\tag{14}$$

According to Cauchy-Schwarz inequality,

$$\begin{aligned} \sum_{k=1}^K \frac{\alpha_k^2}{n_k - 1} \sum_{k=1}^K (n_k - 1) &\geq \left(\sum_{k=1}^K \frac{\alpha_k}{\sqrt{n_k - 1}} \times \sqrt{n_k - 1} \right)^2 \\ &= \left(\sum_{k=1}^K \alpha_k \right)^2 \\ &= 1 \end{aligned} \tag{15}$$

Therefore,

$$\text{var} [\hat{\sigma}^2] \geq \frac{2\sigma^4}{n - K} \tag{16}$$

The equality happens if and only if:

$$\frac{\alpha_1}{n_1 - 1} = \frac{\alpha_2}{n_2 - 1} = \dots = \frac{\alpha_K}{n_K - 1} \tag{17}$$

which can be satisfied by $\alpha_k = (n_k - 1)/(n - K)$, i.e.

$$\frac{\alpha_1}{n_1 - 1} = \frac{\alpha_2}{n_2 - 1} = \dots = \frac{\alpha_K}{n_K - 1} = \frac{1}{n - K} \tag{18}$$

Therefore, $\alpha_k = (n_k - 1)/(n - K)$ minimizes the variance of σ^2 under the Gaussian assumption.

P6

Ans:

Majority vote approach:

X is classified to *Red*: 6 estimates, X is classified to *Green*: 4 estimates. Thus, X is classified to *Red*.

Average probability:

$$\begin{aligned}\bar{\mathbb{P}}[\text{Class is Red}|X] &= \frac{\sum^n \mathbb{P}[\text{Class is Red}|X]}{n} \\ &= \frac{0.1 + 0.15 + 0.2 + 0.2 + 0.55 + 0.6 + 0.6 + 0.65 + 0.7 + 0.75}{10} \\ &= 0.45\end{aligned}$$

Because $\bar{\mathbb{P}}[\text{Class is Red}|X] = 0.45 < 0.5$, X is classified to *Green*.

P7

(a)

Code:

```
1 # (a)
2 rm(list=ls())
3 set.seed(1000)
4 load('/Users/yangchenye/Downloads/OJ.rda')
5 dim(OJ)
6 OJ=na.omit(OJ) # remove incomplete cases
7 dim(OJ)
8 names(OJ)
9
10 # 800: the sample size
11 smp_size = floor(800)
12 train_ind = sample(seq_len(nrow(OJ)), size = smp_size)
13 # select train data and test data
14 train = subset(OJ[train_ind, ])
15 test = subset(OJ[-train_ind, ])
```

(b)

Code:

```
1 # (b)
2 library(tree)
3 tree.fit = tree(train$Purchase ~ ., data=train)
4 summary(tree.fit)
```

Result:

```
1 > summary(tree.fit)
2
3 Classification tree:
4 tree(formula = train$Purchase ~ ., data = train)
5 Variables actually used in tree construction:
6 [1] "LoyalCH"      "PriceDiff"    "SalePriceMM"
7 Number of terminal nodes: 8
8 Residual mean deviance: 0.7486 = 592.9 / 792
9 Misclassification error rate: 0.16 = 128 / 800
```

Ans:

The training error rate is $128/800=0.16$

The tree has 8 terminal nodes.

(c)

Code:

```
1 # (c)
2 tree.fit
```

Result:

```
1 > tree.fit
2 node), split, n, deviance, yval, (yprob)
3      * denotes terminal node
4
5 1) root 800 1066.00 CH ( 0.61500 0.38500 )
6 2) LoyalCH < 0.5036 353 422.60 MM ( 0.28612 0.71388 )
7 4) LoyalCH < 0.276142 170 131.00 MM ( 0.12941 0.87059 )
8 8) LoyalCH < 0.035047 57 10.07 MM ( 0.01754 0.98246 ) *
9 9) LoyalCH > 0.035047 113 108.50 MM ( 0.18584 0.81416 ) *
10 5) LoyalCH > 0.276142 183 250.30 MM ( 0.43169 0.56831 )
11 10) PriceDiff < 0.05 78 79.16 MM ( 0.20513 0.79487 ) *
12 11) PriceDiff > 0.05 105 141.30 CH ( 0.60000 0.40000 ) *
13 3) LoyalCH > 0.5036 447 337.30 CH ( 0.87472 0.12528 )
14 6) LoyalCH < 0.764572 187 206.40 CH ( 0.75936 0.24064 )
15 12) SalePriceMM < 2.125 120 156.60 CH ( 0.64167 0.35833 )
16 24) PriceDiff < -0.35 16 17.99 MM ( 0.25000 0.75000 ) *
17 25) PriceDiff > -0.35 104 126.70 CH ( 0.70192 0.29808 ) *
18 13) SalePriceMM > 2.125 67 17.99 CH ( 0.97015 0.02985 ) *
19 7) LoyalCH > 0.764572 260 91.11 CH ( 0.95769 0.04231 ) *
```

Ans:

For the terminal node:

8) *LoyalCH* < 0.035047 57 10.07 *MM* (0.017540.98246)*

"8)" is the node number.

"*LoyalCH* < 0.035047" is a two-column matrix of the labels for the left and right splits at the node.

"57" is the number of cases reaching that node (*LoyalCH* < 0.035047).

"10.07" is the deviance of the node.

"*MM*" is the fitted value at the node (the mean for regression trees, a majority class for classification trees).

"(0.017540.98246)" is a matrix of fitted probabilities for each response level.

(d)

Code:

```
1 # (d)
2 plot(tree.fit)
3 text(tree.fit, pretty=0)
```

Result:




Figure 1: Regression tree with 8 terminal nodes

Ans:

For customers with $LoyalCH < 0.276142$, they will choose Minute Maid Orange Juice.

For customers with $0.276142 < LoyalCH < 0.5036$ & $PriceDiff < 0.05$, they will choose Minute Maid Orange Juice.

For customers with $0.5036 < LoyalCH < 0.764572$ & $SalePriceMM < 2.125$ & $PriceDiff < -0.35$, they will choose Minute Maid Orange Juice.

Otherwise, customers will choose Citrus Hill Orange Juice.

(e)

Code:

```

1 # (e)
2 tree.pred = predict(tree.fit, test, type="class")
3 with(test, table(tree.pred, Purchase))

```

Result:

```

1 > with(test, table(tree.pred, Purchase))
2      Purchase
3 tree.pred  CH  MM
4      CH 150  38
5      MM  11  71

```

Ans:

The test error rate is $\frac{11+38}{270} = 0.18$

(f)**Code:**

```

1 # (f)
2 cv.fit = cv.tree(tree.fit)
3 cv.fit
4 cv.fit$size[which.min(cv.fit$dev)]

```

Result:

```

1 > cv.fit
2 $size
3 [1] 8 7 6 5 4 3 2 1
4
5 $dev
6 [1] 673.8589 677.7294 675.2341 729.4281 727.2145 772.9214
   773.5642
7 [8] 1067.3829
8
9 $k
10 [1] -Inf 11.87503 12.41171 29.77434 31.80546 39.82936
   41.38321
11 [8] 306.37571
12
13 $method
14 [1] "deviance"
15
16 attr(,"class")
17 [1] "prune" "tree.sequence"
18 > cv.fit$size[which.min(cv.fit$dev)]
19 [1] 8

```

Ans:

The optimal tree size is 8.

(g)**Code:**

```

1 # (g)
2 par(mfrow = c(1, 2))
3 # default plot
4 plot(cv.fit)
5 # better plot
6 plot(cv.fit$size, cv.fit$dev / smp_size, type = "b",
7       xlab = "Tree_Size", ylab = "CV_Misclassification_Rate")

```

Result:




Figure 2: Error rate to tree size

(h)

Code:

```

1 # (h)
2 cv.best = cv.fit$size[which.min(cv.fit$dev / smp_size)] #
   misclassification rate of each tree
3 cv.best

```

Result:

```

1 > cv.best
2 [1] 8

```

Ans:

The tree size corresponds to the lowest cross-validated classification error rate is 8.

(i)

Code:

```

1 # (i)
2 prune.fit = prune.tree(tree.fit, best = 5)
3 plot(prune.fit)
4 text(prune.fit, pretty=0)

```

Result:

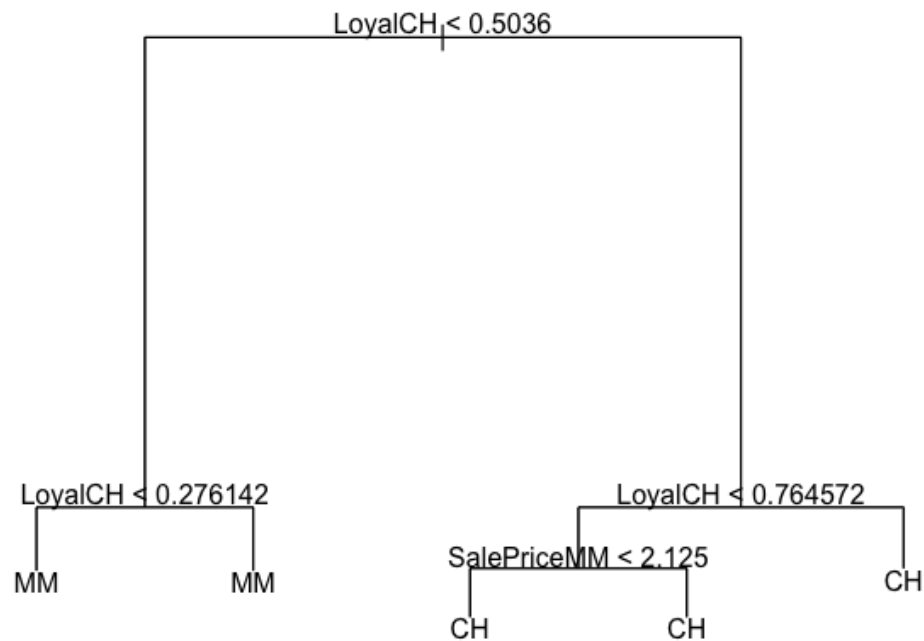


Figure 3: Pruned classification tree

(j)

Code:

```

1 # (j)
2 summary(prune.fit)

```

Result:

```

1 > summary(prune.fit)
2
3 Classification tree:
4 snip.tree(tree = tree.fit, nodes = c(12L, 4L, 5L))
5 Variables actually used in tree construction:
6 [1] "LoyalCH"      "SalePriceMM"
7 Number of terminal nodes: 5
8 Residual mean deviance: 0.8138 = 646.9 / 795
9 Misclassification error rate: 0.1962 = 157 / 800

```

Ans:

Training error rate of unpruned tree: $128/800=0.16$ Training error rate of pruned tree: $157/800=0.1962$

The training error rate of pruned tree is higher.

(k)

Code:

```

1 # (k)

```

```
2 tree.pred = predict(prune.fit, test, type="class")
3 with(test, table(tree.pred, Purchase))
```

Result:

```
1 > with(test, table(tree.pred, Purchase))
2      Purchase
3 tree.pred  CH  MM
4      CH 129  25
5      MM  32  84
```

Ans:

Test error rate of unpruned tree: $(11 + 38)/270 = 0.18$

Test error rate of pruned tree: $(32 + 25)/270 = 0.21$

The test error rate of pruned tree is higher.