

EECS E6690: Statistical Learning for Biological and Information Systems

Lecture 2: Multiple Linear Regression

Prof. Predrag R. Jelenković
Time: Tuesday 4:10-6:40pm
1127 Seeley W. Mudd Building

Dept. of Electrical Engineering
Columbia University , NY 10027, USA
Office: 812 Schapiro Research Bldg.
Phone: (212) 854-8174
Email: predrag@ee.columbia.edu
URL: <http://www.ee.columbia.edu/~predrag>

Last lecture: Intro to stat learning

Supervised vs. Unsupervised learning

Supervised:

- ▶ Let Y be the output variable, and X the input vector

$X = (X_1, X_2, \dots, X_p)$. Then

$$Y = f(X) + \epsilon$$

- ▶ Want to **estimate** f
- ▶ ϵ is unavoidable/irreducible noise that is independent of X , zero mean
- ▶ How to estimate f from the data? How to evaluate the estimate?
- ▶ Errors: irreducible, reducible, bias
- ▶ Overfitting and testing
John von Neumann on overfitting: "With four parameters I can fit an elephant, and with five I can make him wiggle his trunk."

Unsupervised: **No** $f(\cdot)/Y$, just X

Last lecture: Estimation and Testing

Estimation:

- ▶ Select a class of function for f : Hypothesis class \mathcal{H}
say, \mathcal{H} are linear functions, i.e., linear regression
- ▶ Select as distance metric, i.e., **loss function**, which measures the error between $f \in \mathcal{H}$ and data
- ▶ Optimization: find $\hat{f} \in \mathcal{H}$ which minimizes the error/loss function

Testing: How good is \hat{f} on unseen data? Two approaches:

- ▶ Analytical (first 3 lectures)
 - ▶ Make some analytical assumptions, e.g. Gaussian
 - ▶ Compute distributions for the parameters of interest
 - ▶ Develop statistical tests to characterize \hat{f} : t-test, F-test, etc
- ▶ Numerical (rest of the class)
 - ▶ Split data into training and testing
 - ▶ Use training data to find \hat{f}
 - ▶ Use testing data to evaluate how good is \hat{f}

Last lecture

- ▶ Install and get familiar with R (attend the recitation session)

Brief stat review:

- ▶ X_1, X_2, \dots, X_n – i.i.d. with mean μ and variance σ^2
- ▶ Estimators of mean and variance

- ▶ Sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

- ▶ Sample variance

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

- ▶ \bar{X} and S^2 are **unbiased** estimators, i.e.

$$\mathbb{E}\bar{X} = \mu, \quad \text{and} \quad \mathbb{E}S^2 = \sigma^2$$

- ▶ Variability of \bar{X} : $\text{SE}(\bar{X})$ = Standard Error of the mean

$$\text{Var}(\bar{X}) = \sigma^2/n \approx (\text{SE}(\bar{X}))^2 = S^2/n$$

Last lecture: Variability

- ▶ If X_1, \dots, X_n are i.i.d. and **normal/Gaussian**, then
 - ▶ \bar{X} is normal
 - ▶ S^2 has Chi - square distribution:

$$\frac{n-1}{\sigma^2} S^2 \sim \chi_{n-1}^2,$$

χ_{n-1}^2 = sum of $(n-1)$ squares of independent standard normal variables

- ▶ \bar{X} and S^2 are independent
- ▶ t -value and Student's t -distribution:

$$\frac{\bar{X} - \mu}{\sqrt{S^2/n}} \sim \frac{\mathcal{N}(0, 1)}{\sqrt{\chi_{n-1}^2/(n-1)}} = t_{n-1},$$

William Gosset, 1908, under pen name Student

- ▶ ... if X is not normal/Gaussian, then use the CLT

Last lecture: Hypothesis testing - t -test

t_n has a known symmetric and bell shaped density
(use $\Gamma(k + 1/2) \approx \sqrt{k}\Gamma(k)$ for large k)

$$f_n(t) = \frac{\Gamma((n+1)/2)}{\sqrt{\pi n} \Gamma(n/2)} \left(1 + \frac{t^2}{n}\right)^{-\frac{n+1}{2}} \approx \frac{1}{\sqrt{2\pi}} e^{-\frac{t^2}{2}} \quad (\text{large } n)$$

t -test:

- ▶ Null hypothesis $\mathcal{H}_0 : \mu = \mu_0$
- ▶ Under \mathcal{H}_0 , compute t -value and p -value:

$$t = \frac{\bar{X} - \mu_0}{\sqrt{S^2/n}}, \quad p = \mathbb{P}[|t_{n-1}| \geq |t|]$$

- ▶ Since large values of t unlikely under \mathcal{H}_0 , typically
 - ▶ pick a significance value, say $\alpha = 0.05$
 - ▶ reject \mathcal{H}_0 if $p < \alpha$, say $p < 0.05$
 - ▶ accept \mathcal{H}_0 if $p \geq \alpha$, say $p \geq 0.05$

Last lecture: Linear regression

- ▶ Simple approach to supervised learning
- ▶ Assumes linear dependence of Y on X_1, X_2, \dots, X_p
Almost never true in reality.
- ▶ Extremely useful both conceptually and practically
- ▶ Linear model

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- ▶ Estimate β_0 and β_1 by **minimizing residuals**

$$y_i - \hat{y}_i = y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)$$

- ▶ Norm selection (distance measure) is important
e.g., l_2 vs. l_1
- ▶ l_2 regression: Least squares (r_{xy} - correlation coefficient)

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})}{\sum_{i=1}^n (x_i - \bar{x})^2} = r_{xy} \frac{S_y}{S_x}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

Last lecture: Statistics of $\hat{\beta}_0$ and $\hat{\beta}_1$

- ▶ **Repeated sampling**
- ▶ $\hat{\beta}_0$ and $\hat{\beta}_1$ vary
- ▶ **Unbiased estimators:**

$$\mathbb{E}\hat{\beta}_0 = \beta_0 \quad \text{and} \quad \mathbb{E}\hat{\beta}_1 = \beta_1$$

- ▶ Variances: (model $Y = f(X) + \epsilon$, ϵ -Gaussian)

$$\text{Var}(\hat{\beta}_1) = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2},$$

$$\text{Var}(\hat{\beta}_0) = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right),$$

where $\sigma^2 = \text{Var}(\epsilon)$

- ▶ An estimate of σ^2 :

$$\text{RSE}^2 = \frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1}{n-2} \text{RSS},$$

where RSE is the Residual Standard Error

Last lecture: Hypothesis testing and confidence intervals

- ▶ Normality assumption: $\epsilon \sim \mathcal{N}(0, \sigma^2)$

- ▶ t -statistic:

$$\frac{\hat{\beta}_1 - \beta_1}{\text{SE}(\hat{\beta}_1)} \sim t_{n-2},$$

where

$$\text{SE}(\hat{\beta}_1)^2 = \frac{1}{n-2} \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

- ▶ Hypothesis testing using t statistics
- ▶ $(1 - \gamma)$ confidence interval (say $\gamma = 5\%$, $1 - \gamma = 95\%$):

$$[\hat{\beta}_1 - \text{SE}(\hat{\beta}_1) \cdot t_{\gamma/2, n-2}, \hat{\beta}_1 + \text{SE}(\hat{\beta}_1) \cdot t_{\gamma/2, n-2}]$$

where $t_{\gamma/2, n-2}$ is the $(1 - \gamma/2)$ -th quantile of the t_{n-2} distribution $\mathbb{P}[-t_{\gamma/2, n-2} \leq t_{n-2} \leq t_{\gamma/2, n-2}] = 1 - \gamma$, i.e.,

$$\mathbb{P}[\hat{\beta}_1 - \text{SE}(\hat{\beta}_1) \cdot t_{\gamma/2, n-2} \leq \beta_1 \leq \hat{\beta}_1 + \text{SE}(\hat{\beta}_1) \cdot t_{\gamma/2, n-2}] = 1 - \gamma$$

Multiple linear regression

- ▶ Model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

- ▶ Example

$$\text{Sales} = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio} + \beta_3 \times \text{Newspaper} + \epsilon$$

- ▶ Interpretation: β_i is the average effect on Y of a one unit increase in X_i , holding all other predictors fixed

- ▶ Notes

- ▶ Ideally the predictors are uncorrelated
- ▶ Correlations amongst predictors cause problems
 - ▶ increased variance of coefficients
 - ▶ tricky interpretations (example: $X_1 = X_2^2$)
- ▶ Claims of causality should be avoided for observational data

l_2 regression

- ▶ n observations: $(y_i, x_{i,1}, x_{i,2}, \dots, x_{i,p})$
- ▶ Given estimates $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$, the prediction is

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \dots + \hat{\beta}_p x_p,$$

or in matrix form

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_n \end{bmatrix} = \hat{\mathbf{y}} = \mathbf{X} \hat{\boldsymbol{\beta}} = \begin{bmatrix} 1 & x_{1,1} & \cdots & x_{1,p} \\ 1 & x_{2,1} & \cdots & x_{2,p} \\ \vdots & & & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,p} \end{bmatrix} \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \vdots \\ \hat{\beta}_p \end{bmatrix}$$

- ▶ Minimize (over β_1, \dots, β_p) the residual sum of squares

$$\text{RSS}(\boldsymbol{\beta}) = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

l_2 regression: Solution


- ▶ Minimize RSS: Differentiating $RSS(\beta)$, we get

$$\frac{\partial RSS(\beta)}{\partial \beta} = -2\mathbf{X}^\top(\mathbf{y} - \mathbf{X}\beta) = \mathbf{0}$$

- ▶ Solution $\beta = (\hat{\beta}_1, \dots, \hat{\beta}_p)$: assuming $\mathbf{X}^\top \mathbf{X}$ is full rank

$$\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \quad \hat{\mathbf{y}} = \mathbf{X} \hat{\beta}$$

- ▶ Geometry



Example: Advertising data

```
> lm2<-lm(adv$Sales~adv$TV+adv$Radio+adv$Newspaper)
> summary(lm2)
```

Call:

```
lm(formula = adv$Sales ~ adv$TV + adv$Radio + adv$Newspaper)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8.8277	-0.8908	0.2418	1.1893	2.8292

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.938889	0.311908	9.422	<2e-16 ***
adv\$TV	0.045765	0.001395	32.809	<2e-16 ***
adv\$Radio	0.188530	0.008611	21.893	<2e-16 ***
adv\$Newspaper	-0.001037	0.005871	-0.177	0.86

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.686 on 196 degrees of freedom

Multiple R-squared: 0.8972, Adjusted R-squared: 0.8956

F-statistic: 570.3 on 3 and 196 DF, p-value: < 2.2e-16

```
> cor(adv[,2:5])
```

	TV	Radio	Newspaper	Sales
TV	1.00000000	0.05480866	0.05664787	0.7822244
Radio	0.05480866	1.00000000	0.35410375	0.5762226
Newspaper	0.05664787	0.35410375	1.00000000	0.2282990
Sales	0.78222442	0.57622257	0.22829903	1.0000000

Solution: Algebraic/geometric interpretations

- ▶ Ideally $\mathbf{y} = \mathbf{X}\hat{\boldsymbol{\beta}}$ (RSS = 0), but this equation has no solution (except in trivial cases), since $\mathbf{y} \notin C(\mathbf{X})$
 $C(\mathbf{X})$ = column space, i.e., hyperplane formed by columns of \mathbf{X} .
- ▶ Instead, we solve $\mathbf{P}\mathbf{y} = \mathbf{X}\hat{\boldsymbol{\beta}}$, where $\mathbf{P} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ is the l_2 -projection matrix onto $C(\mathbf{X})$
 - ▶ \mathbf{P} is sometimes called "hat" matrix, denoted as \mathbf{H} , since it puts a hat on \mathbf{y} , i.e. $\hat{\mathbf{y}} = \mathbf{P}\mathbf{y} \equiv \mathbf{H}\mathbf{y}$
- ▶ Equivalently, $\hat{\boldsymbol{\beta}}$ satisfies

$$\mathbf{X}^\top \mathbf{y} = \mathbf{X}^\top \mathbf{X} \hat{\boldsymbol{\beta}}$$

- ▶ A unique solution exists when the columns of \mathbf{X} are linearly independent – in that case, $\mathbf{X}^\top \mathbf{X}$ is full-rank and positive definite
- ▶ Consequences:
 - ▶ $(\mathbf{y} - \hat{\mathbf{y}}) = (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}})$ is perpendicular to $C(\mathbf{X})$
 - ▶ $\mathbf{0} = (\mathbf{y} - \hat{\mathbf{y}})^\top \mathbf{X} = (\mathbf{y} - \mathbf{P}\mathbf{y})^\top \mathbf{X} = \mathbf{y}^\top (\mathbf{X} - \mathbf{P}\mathbf{X})$
 - ▶ $\sum_{i=1}^n (y_i - \hat{y}_i) = (\mathbf{y} - \hat{\mathbf{y}})^\top \mathbf{1} = 0$
 - ▶ $\|\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}\|_2^2 = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + (\hat{\boldsymbol{\beta}} - \boldsymbol{\beta})^\top \mathbf{X}^\top \mathbf{X} (\hat{\boldsymbol{\beta}} - \boldsymbol{\beta})$

How good is the model fit?

- ▶ Total sum of squares: $TSS = (\mathbf{y} - \bar{y}\mathbf{1})^\top (\mathbf{y} - \bar{y}\mathbf{1})$
- ▶ Explained sum of squares: $ESS = (\hat{\mathbf{y}} - \bar{y}\mathbf{1})^\top (\hat{\mathbf{y}} - \bar{y}\mathbf{1})$
- ▶ Then

$$\begin{aligned} TSS &= (\mathbf{y} - \hat{\mathbf{y}} + \hat{\mathbf{y}} - \bar{y}\mathbf{1})^\top (\mathbf{y} - \hat{\mathbf{y}} + \hat{\mathbf{y}} - \bar{y}\mathbf{1}) \\ &= RSS + ESS + 2(\mathbf{y} - \hat{\mathbf{y}})^\top (\hat{\mathbf{y}} - \bar{y}\mathbf{1}) \\ &= RSS + ESS + 2(\mathbf{y} - \hat{\mathbf{y}})^\top (\mathbf{X}\hat{\boldsymbol{\beta}} - \mathbf{1}\bar{y}) \\ &= RSS + ESS \end{aligned}$$

- ▶ A measure of quality of the model

$$R^2 = \frac{ESS}{TSS} = 1 - \frac{RSS}{TSS}$$

- ▶ $R^2 \uparrow$ as more explanatory variables are added to the model – need to consider the number of variables

Example: Advertising data

```
> summary(lm(adv$Sales~adv$TV+adv$Radio))
```

Call:

```
lm(formula = adv$Sales ~ adv$TV + adv$Radio)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8.7977	-0.8752	0.2422	1.1708	2.8328

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.92110	0.29449	9.919	<2e-16 ***
adv\$TV	0.04575	0.00139	32.909	<2e-16 ***
adv\$Radio	0.18799	0.00804	23.382	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.681 on 197 degrees of freedom

Multiple R-squared: 0.8972, Adjusted R-squared: 0.8962

F-statistic: 859.6 on 2 and 197 DF, p-value: < 2.2e-16

► R^2

► Are all predictors useful? Which are?

Distribution of $\hat{\beta}$

- ▶ Normality assumption: i.i.d. $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$
- ▶ $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$ is also normally distributed, with mean $\boldsymbol{\mu} = \mathbf{X}\boldsymbol{\beta}$, covariance matrix $\boldsymbol{\Sigma} = \sigma^2 \mathbf{I}$ and density

$$f_{\mathbf{y}}(\mathbf{x}) = \frac{1}{(2\pi)^{p/2} |\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})}$$

- ▶ $\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ is also normal with

$$\mathbb{E}\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{X} \boldsymbol{\beta} + (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbb{E}\boldsymbol{\epsilon} = \boldsymbol{\beta}$$

and

$$\begin{aligned} \text{Cov}(\hat{\boldsymbol{\beta}}) &= \mathbb{E}(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta})(\hat{\boldsymbol{\beta}} - \boldsymbol{\beta})^\top \\ &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \boldsymbol{\epsilon} ((\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \boldsymbol{\epsilon})^\top = \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1} \end{aligned}$$

- ▶ Hence $\hat{\boldsymbol{\beta}} \sim \mathcal{N}(\boldsymbol{\beta}, \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1})$

Residuals

- Residuals $\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}$ are also normal with

$$\mathbb{E}[\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}] = \mathbb{E}[\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon} - \mathbf{X}\hat{\boldsymbol{\beta}}] = \mathbf{0}$$

and

$$\begin{aligned}\mathbb{E}(\mathbf{y} - \hat{\mathbf{y}})(\mathbf{y} - \hat{\mathbf{y}})^\top &= \mathbb{E}(\mathbf{y} - \mathbf{P}\mathbf{y})(\mathbf{y} - \mathbf{P}\mathbf{y})^\top \\ &= \mathbb{E}(\mathbf{I} - \mathbf{P})\boldsymbol{\epsilon}\boldsymbol{\epsilon}^\top(\mathbf{I} - \mathbf{P})^\top = \sigma^2(\mathbf{I} - \mathbf{P})\end{aligned}$$

since

$$\mathbf{y} - \hat{\mathbf{y}} = (\mathbf{I} - \mathbf{P})\mathbf{y} = (\mathbf{I} - \mathbf{P})(\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}) = (\mathbf{I} - \mathbf{P})\boldsymbol{\epsilon}$$

Estimating σ

- ▶ $\text{RSS} = (\mathbf{y} - \hat{\mathbf{y}})^\top (\mathbf{y} - \hat{\mathbf{y}}) = \boldsymbol{\epsilon}^\top (\mathbf{I} - \mathbf{P}) \boldsymbol{\epsilon}$, where

$$\begin{aligned}\text{tr}(\mathbf{I} - \mathbf{P}) &= \text{tr}(\mathbf{I} - \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top) \\ &= n - \text{tr}(\mathbf{X}^\top \mathbf{X} (\mathbf{X}^\top \mathbf{X})^{-1}) = n - p - 1\end{aligned}$$

- ▶ Then it can be shown

$$\frac{\text{RSS}}{\sigma^2} \sim \chi_{n-p-1}^2$$

- ▶ Estimator:

$$\hat{\sigma} = \sqrt{\frac{\text{RSS}}{n - p - 1}}$$

Back to testing

- ▶ $\mathcal{H}_0 : \beta_j = 0$
- ▶ Intuition: reject \mathcal{H}_0 if $\hat{\beta}_j$ is “large”
- ▶ How large?
- ▶ Under \mathcal{H}_0 , $\hat{\beta}_j$ is $\mathcal{N}(0, \sigma^2(\mathbf{X}^\top \mathbf{X})_{j,j}^{-1})$
- ▶ Consider t -statistic

$$\frac{\hat{\beta}_j}{\hat{\sigma} \sqrt{(\mathbf{X}^\top \mathbf{X})_{j,j}^{-1}}} = \frac{\frac{\hat{\beta}_j}{\sigma \sqrt{(\mathbf{X}^\top \mathbf{X})_{j,j}^{-1}}}}{\sqrt{\frac{\text{RSS}}{\sigma^2(n-p-1)}}} \sim t_{n-p-1}$$

F-test

- ▶ Better idea: use RSS to test instead of $\hat{\beta}$
- ▶ $\mathcal{H}_0 : \beta_1 = \beta_2 = \dots = \beta_p = 0$
- ▶ \mathcal{H}_1 : exists j such that $\beta_j \neq 0$
- ▶ Under \mathcal{H}_0 , we have a null model: $Y = \beta_0 + \epsilon$
- ▶ Let RSS_0 be the residual sum of squares under \mathcal{H}_0
- ▶ Under \mathcal{H}_0 :

$$\frac{\text{RSS}_0 - \text{RSS}}{\sigma^2} = \frac{\text{TSS} - \text{RSS}}{\sigma^2} \sim \chi_p^2$$

and

$$\frac{\frac{\text{TSS} - \text{RSS}}{p}}{\frac{\text{RSS}}{n-p-1}} \sim F_{p, n-p-1}$$

- ▶ Back to the example:

Residual standard error: 1.686 on 196 degrees of freedom
Multiple R-squared: 0.8972, Adjusted R-squared: 0.8956
F-statistic: 570.3 on 3 and 196 DF, p-value: < 2.2e-16

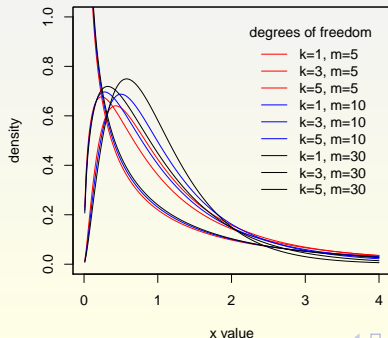
F-distribution

- ▶ $F_{k,m}$: independent $V \sim \chi_k^2$ and $W \sim \chi_m^2$

$$\frac{V/k}{W/m} \sim F_{k,m}$$

- ▶ Mean ($m > 2$): $m/(m-2)$
- ▶ Variance ($m > 4$): $\frac{2m^2(k+m-2)}{k(m-2)^2(m-4)}$

PDFs of F distributions



F-test

- ▶ $\mathcal{H}_0 : \beta_j = 0$
- ▶ Under \mathcal{H}_0 , we have a reduced model:

$$Y = \beta_0 + \beta_1 X_1 + \cdots + \beta_{j-1} X_{j-1} + \beta_{j+1} X_{j+1} + \cdots + \beta_p X_p + \epsilon$$

- ▶ Refer to the reduced model by index $-j$
- ▶ Intuition: while $RSS \leq RSS_{-j}$...

 - ▶ if $RSS \ll RSS_{-j}$, then reject \mathcal{H}_0
 - ▶ if $RSS \approx RSS_{-j}$, then accept \mathcal{H}_0

- ▶ Under \mathcal{H}_0 :

$$\frac{RSS_{-j} - RSS}{\sigma^2} \sim \chi_1^2$$

and

$$\frac{\frac{RSS_{-j} - RSS}{RSS}}{\frac{1}{n-p-1}} \sim F_{1,n-p-1}$$

F-test

- ▶ (m) denotes a sub-model obtained by a linear constraint on β
- ▶ Examples
 - ▶ $\beta_1 = \beta_2 = \dots = \beta_p$: $Y = \beta_0 + \epsilon$
 - ▶ $\beta_1 = \beta_2$: $Y = \beta_0 + \beta_1(X_1 + X_2) + \beta_3X_3 + \dots + \beta_pX_p + \epsilon$
- ▶ Testing: \mathcal{H}_0 (reduced model) vs. \mathcal{H}_1 (complete model)
- ▶ $q < p$ is the number of explanatory variables in the reduced model
- ▶ Under \mathcal{H}_0 :

$$\frac{\text{RSS}_{(m)} - \text{RSS}}{\sigma^2} \sim \chi_{p-q}^2$$

Qualitative predictors

- ▶ Credit.csv data set
- ▶ 400 observations:

```
"", "Income", "Limit", "Rating", "Cards", "Age", "Education", "Gender", "Student", "Married", "Ethnicity", "Balance"  
"1", 14.891, 3606, 283, 2, 34, 11, " Male", "No", "Yes", "Caucasian", 333  
"2", 106.025, 6645, 483, 3, 82, 15, "Female", "Yes", "Yes", "Asian", 903  
"3", 104.593, 7075, 514, 4, 71, 11, " Male", "No", "No", "Asian", 580  
"4", 148.924, 9504, 681, 3, 36, 11, "Female", "No", "No", "Asian", 964  
"5", 55.882, 4897, 357, 2, 68, 16, " Male", "No", "Yes", "Caucasian", 331  
.  
.  
.  
"398", 57.872, 4171, 321, 5, 67, 12, "Female", "No", "Yes", "Caucasian", 138  
"399", 37.728, 2525, 192, 1, 44, 13, " Male", "No", "Yes", "Caucasian", 0  
"400", 18.701, 5524, 415, 5, 64, 7, "Female", "No", "No", "Asian", 966
```

- ▶ Quantitative predictors: Income (in thousands), Limit (credit), Rating, Cards (number of), Age, Education (years of), Balance
- ▶ Qualitative predictors (factors): Gender, Student, Married, Ethnicity

Qualitative predictors

- ▶ Dependency of Balance on Gender
- ▶ Ignore all other variables
- ▶ Gender has two levels:

$$X_i = \begin{cases} 1, & \text{if } i\text{th individual is female} \\ 0, & \text{if } i\text{th individual is male} \end{cases}$$

- ▶ Model: $Y = \beta_0 + \beta_1 X + \epsilon$
- ▶ Interpretation
 - ▶ β_0 : average Balance among males
 - ▶ $\beta_0 + \beta_1$: average Balance among females
 - ▶ β_1 : average difference in Balance between females and males
- ▶ The 1/0 encoding is arbitrary. Can use another scheme – only the interpretation changes

Example

```
> credit <- read.csv("credit.csv",header=TRUE,sep=",")
> lm3<-lm(Balance~Gender,data=credit)
> summary(lm3)
```

Call:

```
lm(formula = Balance ~ Gender, data = credit)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-529.54	-455.35	-60.17	334.71	1489.20

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	509.80	33.13	15.389	<2e-16 ***
GenderFemale	19.73	46.05	0.429	0.669

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 460.2 on 398 degrees of freedom

Multiple R-squared: 0.0004611, Adjusted R-squared: -0.00205

F-statistic: 0.1836 on 1 and 398 DF, p-value: 0.6685

Qualitative predictors

- ▶ When a factor has more than two levels, a single dummy variable can not represent all possible values
- ▶ In that case, create additional dummy variables
- ▶ Example: Ethnicity

$$X_{i,1} = \begin{cases} 1, & \text{if } i\text{th individual is Asian} \\ 0, & \text{if } i\text{th individual is not Asian} \end{cases}$$

$$X_{i,2} = \begin{cases} 1, & \text{if } i\text{th individual is Caucasian} \\ 0, & \text{if } i\text{th individual is not Caucasian} \end{cases}$$

- ▶ Model:

$$Y = \begin{cases} \beta_0 + \beta_1 + \epsilon, & \text{if } i\text{th individual is Asian} \\ \beta_0 + \beta_2 + \epsilon, & \text{if } i\text{th individual is Caucasian} \\ \beta_0 + \epsilon, & \text{otherwise} \end{cases}$$

Example

```
> lm4<-lm(Balance~Ethnicity,data=credit)
> summary(lm4)
```

Call:

```
lm(formula = Balance ~ Ethnicity, data = credit)
```

Residuals:

Min	1Q	Median	3Q	Max
-531.00	-457.08	-63.25	339.25	1480.50

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	531.00	46.32	11.464	<2e-16 ***
EthnicityAsian	-18.69	65.02	-0.287	0.774
EthnicityCaucasian	-12.50	56.68	-0.221	0.826

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 460.9 on 397 degrees of freedom

Multiple R-squared: 0.0002188, Adjusted R-squared: -0.004818

F-statistic: 0.04344 on 2 and 397 DF, p-value: 0.9575

Extensions: Interactions

- ▶ Relax additivity
- ▶ Back to Advertising data set
- ▶ Suppose spending money on radio advertising increases the effectiveness of TV advertising
- ▶ New model:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2 + \epsilon$$

- ▶ $X_1 X_2$ – just multiply observations
- ▶ Hierarchical principle: if interaction is in the model, main effects are in the model, even if main effects are not significant

Example

```
> lm5<-lm(Sales~TV+Radio+TV*Radio,data=adv)
> summary(lm5)

Call:
lm(formula = Sales ~ TV + Radio + TV * Radio, data = adv)


Residuals:
    Min       1Q   Median       3Q      Max
-6.3366 -0.4028  0.1831  0.5948  1.5246

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  6.750e+00  2.479e-01  27.233  <2e-16 ***
TV           1.910e-02  1.504e-03  12.699  <2e-16 ***
Radio        2.886e-02  8.905e-03   3.241  0.0014 **
TV:Radio      1.086e-03  5.242e-05  20.727  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9435 on 196 degrees of freedom
Multiple R-squared:  0.9678, Adjusted R-squared:  0.9673
F-statistic: 1963 on 3 and 196 DF, p-value: < 2.2e-16
```

- ▶ $\hat{\beta}_3$: the increase in the effectiveness of TV advertising for a one unit increase in radio advertising (or vice-versa)

Example




Example

- ▶ Credit data set
- ▶ $Y = \text{Balance}$, $X_1 = \text{Income}$, $X_2 = \text{Student} \in \{0, 1\}$
- ▶ Two models:

$$M_1 : Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$$

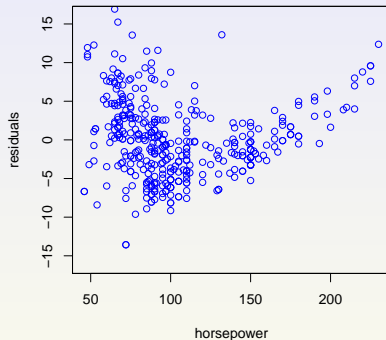

$$M_2 : Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2 + \epsilon$$



- ▶ Changes in income affect students and non-students differently



Extensions: Nonlinearities - Basis expansion

- ▶ Auto data set: mpg vs. horsepower



- ▶ Polynomial regression
- ▶ Model: $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \epsilon$
- ▶ The model is still linear in β : treat X^2 as a variable



Example



- ▶ Increasing the degree can cause overfitting
- ▶ Alternative transformations

Colinearity

- ▶ Credit data set:



- ▶ Rating and Limit are co-linear
- ▶ Difficult to assess individual impact on Balance
- ▶ $\mathbf{X}^\top \mathbf{y} = \mathbf{X}^\top \mathbf{X} \hat{\beta}$ can be numerically unstable



Example

```
> lm8a<-lm(Balance ~ Limit + Age,data=credit)
> summary(lm8a)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-173.410901	43.828387048	-3.956589	9.005366e-05
Limit	0.173365	0.005025662	34.495944	1.627198e-121
Age	-2.291486	0.672484540	-3.407492	7.226468e-04

```
> lm8b<-lm(Balance ~ Limit + Rating,data=credit)
> summary(lm8b)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-377.53679536	45.25417619	-8.3425846	1.213565e-15
Limit	0.02451438	0.06383456	0.3840298	7.011619e-01
Rating	2.20167217	0.95229387	2.3119672	2.129053e-02



Prediction considerations

- ▶ Assumptions
- ▶ Uncertainties/Errors
 - ▶ Regression coefficients are noisy
 - ▶ The true function might not be linear (bias)
 - ▶ Measurements are noisy even when the function is known

K-NN regression: Non-parametric approach

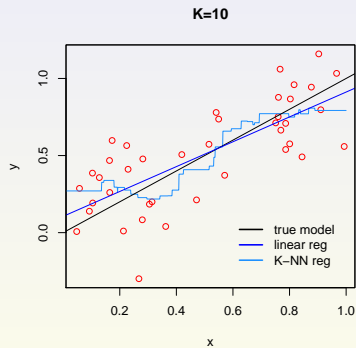

- ▶ Linear regression is not the only approach
- ▶ K-NN: K-nearest neighbors
- ▶ Intuition: “similar” argument values should lead to “similar” function values
 - ▶ distances
 - ▶ data normalization
 - ▶ high dimensionality case
- ▶ Let \mathcal{N}_x^K be the K -nearest neighbors set of observations for x :

$$\hat{f}(x) = \frac{1}{K} \sum_{i \in \mathcal{N}_x^K} y_i$$

- ▶ K is a parameter of the algorithm
 - ▶ Small K – flexible fit, high variance
 - ▶ Large K – smooth, high bias
 - ▶ How to select K ?

Examples

► Linear model:



Reading:

ISL: Finish reading Chapter 3

ESL: Chapter 3

Homework: to be emailed separately.